




Transformers in Vision

AN INTRODUCTION

By:

Dr. Safouane El Ghazouali

BIO



Safouane El Ghazouali, Ph. D.

- Senior researcher in AI at TOELT
- Senior data scientist in Computer Vision
- Experienced in VLM / LLM inference and fine-tuning
- Built a web app UI for fast annotation (VisioFirm)
- Experienced in training / fine-tuning of ML model for domain-specific application

Background

- Ph.D. in AI applied to 3D metrology, LNE, Paris
- Master in computer vision and embedded systems, Polytechnique Haut-de-France, Valenciennes

GitHub repo for the course :

<https://github.com/safouaneelg/HSLU-Transformers-in-Vision>


Transformers in Vision Course - HSLU

Welcome to the GitHub repository for the "Transformers in Vision" course at HSLU.

This repo contains materials and hands-on exercises for Day 3 and Day 4, from the bootcamp starting August 25, 2026.

HSLU Hochschule Luzern Transformers in vision (ViTs)

OVERVIEW OF THE ViT-ARCHITECTURE:



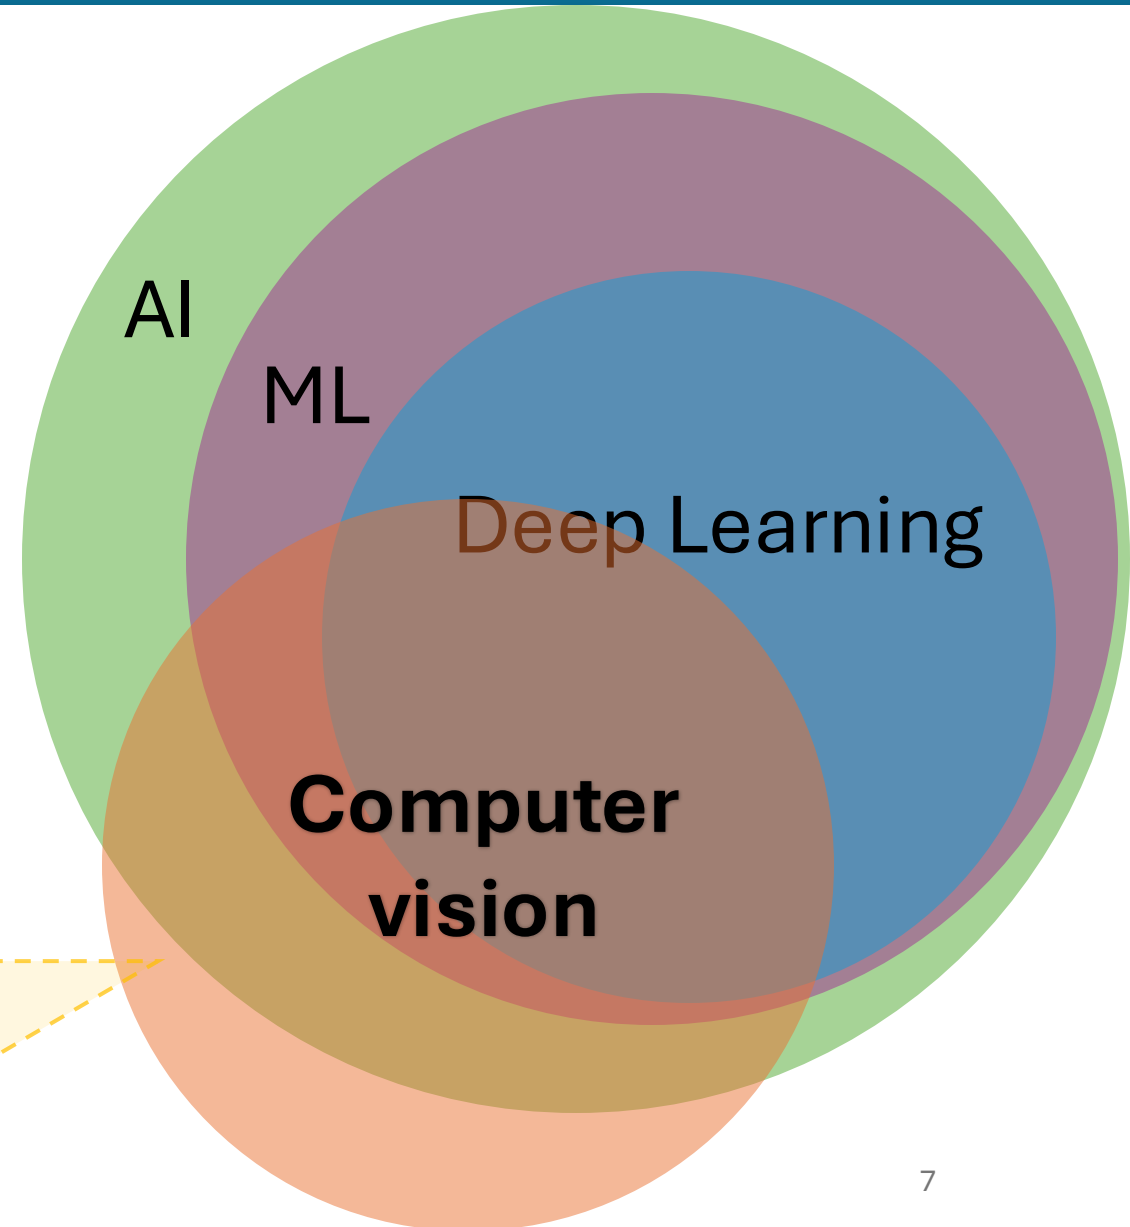
- General Introduction
- Overview of neural architectures evolution
- Transformers in vision (ViTs)
- Divers ViTs architectures
- Hands-on (using “Torch” and ”timm” frameworks):
 - Pre-trained classification model
 - Visualization of features map
 - Finetuning ViT on classification task
- ViTs in object detection
- Hands-on Ultralytics (YOLO models)
 - Inferencing pre-trained model
 - Finetuning YOLOv10 on custom dataset
- Huggingface Transformers library
- Limitations of ViTs

- Background & Motivations
- Applications: image captioning, retrieval, grounding, generative tasks
- CLIP and Zero-Shot Learning
- Hands-on:
 - Inference CLIP for zero-shot classification
 - GroundingDINO for zero-shot detection
 - Prompt Engineering
- BLIP and Multimodal Generation
- Diffusion Model – introduction to Markov Chain
- Hands-on:
 - Stable-diffusion
 - Prompt engineering

General introduction

Computer vision includes:

- 1- Image/signal processing
- 2- 3D reconstruction
- 3- Feature extraction
- 4- Image restoration
- 5- Image classification
- 6- Object detection
- 7- Semantic/instance segmentation
- 8- Scene understanding
- 9- Image generation
- 10- Optical flow analysis
- 11- Pose estimation
- 12- Video tracking



Overview of neural architectures evolution

Computer vision evolution timeline

1950	1970	2017	2020
Foundations in optics and early mathematical-based processing. Focus on understanding light, vision, pinhole and basic digital imaging.	Shift to more structured NNs with hierarchical feature learning. Introduction of convolutional ideas for shift-invariant pattern recognition.	<p>Emergence of advanced modular architectures and powerful backbones able to extract deep features.</p> <p>The most known backbones</p> <ul style="list-style-type: none">- VGGNet 16–19 layers with uniform 3x3 convolutions for depth exploration- AlexNet: 8-layer CNN winning ImageNet challenge- DenseNet: Layer-wise connections for feature reuse	<p>Rise of pure vision transformers, outperforming CNNs on large datasets. Multimodal integration begins.</p> <p>Vision Transformer: Patched images fed into transformers for classification.</p> <p>DeiT (2021): Data-efficient ViT variant.</p> <p>Swin Transformer (2021): Hierarchical shifted windows for efficiency.</p>

Transformers in vision (ViTs)

Original Tranformer paper

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

First ViT paper

**AN IMAGE IS WORTH 16X16 WORDS:
TRANSFORMERS FOR IMAGE RECOGNITION AT SCALE**

**Alexey Dosovitskiy^{*,†}, Lucas Beyer^{*}, Alexander Kolesnikov^{*}, Dirk Weissenborn^{*},
Xiaohua Zhai^{*}, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer,
Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, Neil Houlsby^{*,†}**

^{*}equal technical contribution, [†]equal advising

Google Research, Brain Team

{adosovitskiy, neilhoulby}@google.com

Objectives:

- 1. Demonstrate the Applicability of Pure Transformers to Image Recognition:** Show that a standard Transformer architecture can be applied directly to sequences of image patches for classification tasks, without relying on convolutional neural networks (CNNs), thereby challenging CNN dominance in computer vision.
- 2. Evaluate Performance on Large-Scale Pre-Training and Transfer Learning:** Assess the Vision Transformer (ViT) model's performance when pre-trained on massive datasets (e.g., ImageNet-21k or JFT-300M) and fine-tuned on various benchmarks (e.g., ImageNet, CIFAR-100, VTAB), comparing it to state-of-the-art CNNs.
- 3. Investigate the Impact of Dataset Scale on Transformer Performance:** Explore how training dataset size influences ViT's effectiveness, emphasizing that large-scale data can compensate for Transformers' lack of built-in inductive biases like those in CNNs (e.g., translation equivariance and locality).
- 4. Compare Computational Efficiency and Resource Requirements:** Highlight that ViT achieves strong results with significantly lower computational costs during training compared to leading CNNs, using metrics such as TPUv3-core-days.

Aspect	NLP Tokenization	Vision Tokenization
Definition	Process of breaking down text into smaller units (tokens) such as words, subwords, or characters.	Process of dividing an image into fixed-size patches or regions to process as tokens.
Normalization	<ul style="list-style-type: none"> - Converts text to a standard format (e.g., lowercasing, removing accents). - Removes punctuation or special characters. - Handles contractions (e.g., "don't" → "do not"). 	<ul style="list-style-type: none"> - Normalizes pixel values (e.g., scaling to [0,1] or standardizing with mean and std). - Adjusts color channels (e.g., RGB normalization). - Resizes images to a fixed resolution.
Segmentation/Splitting	<ul style="list-style-type: none"> - Splits text into tokens based on whitespace, punctuation, or subword units (e.g., WordPiece, BPE). - Example: "I am running" → ["I", "am", "running"]. - Subword tokenization for rare words (e.g., "unhappiness" → ["un", "##happiness"]). 	<ul style="list-style-type: none"> - Divides image into fixed-size patches (e.g., 16x16 pixels). - Each patch is treated as a token. - Example: A 224x224 image with 16x16 patches → 196 patches (tokens).
Cleanup	<ul style="list-style-type: none"> - Removes irrelevant tokens (e.g., extra whitespace, stop words like "the"). - Filters out low-frequency tokens or replaces them with [UNK]. - Handles encoding issues (e.g., UTF-8 normalization). 	<ul style="list-style-type: none"> - Removes or adjusts noisy patches (e.g., handling corrupted pixels). - Applies data augmentation (e.g., random cropping, flipping) to enhance robustness. - Filters out irrelevant regions (e.g., background suppression in some cases).
Token Representation	<ul style="list-style-type: none"> - Tokens are mapped to vocabulary indices or embeddings (e.g., Word2Vec, BERT embeddings). - Each token is a discrete unit in a vocabulary. 	<ul style="list-style-type: none"> - Patches are flattened and projected to a fixed-dimensional embedding space (e.g., linear projection in ViT). - Each patch is a vector in a continuous space.
Handling Variability	<ul style="list-style-type: none"> - Handles linguistic variations (e.g., synonyms, misspellings, multilingual text). - Uses subword tokenization to manage out-of-vocabulary words. 	<ul style="list-style-type: none"> - Handles variations in lighting, orientation


OVERVIEW OF THE ViT-ARCHITECTURE:



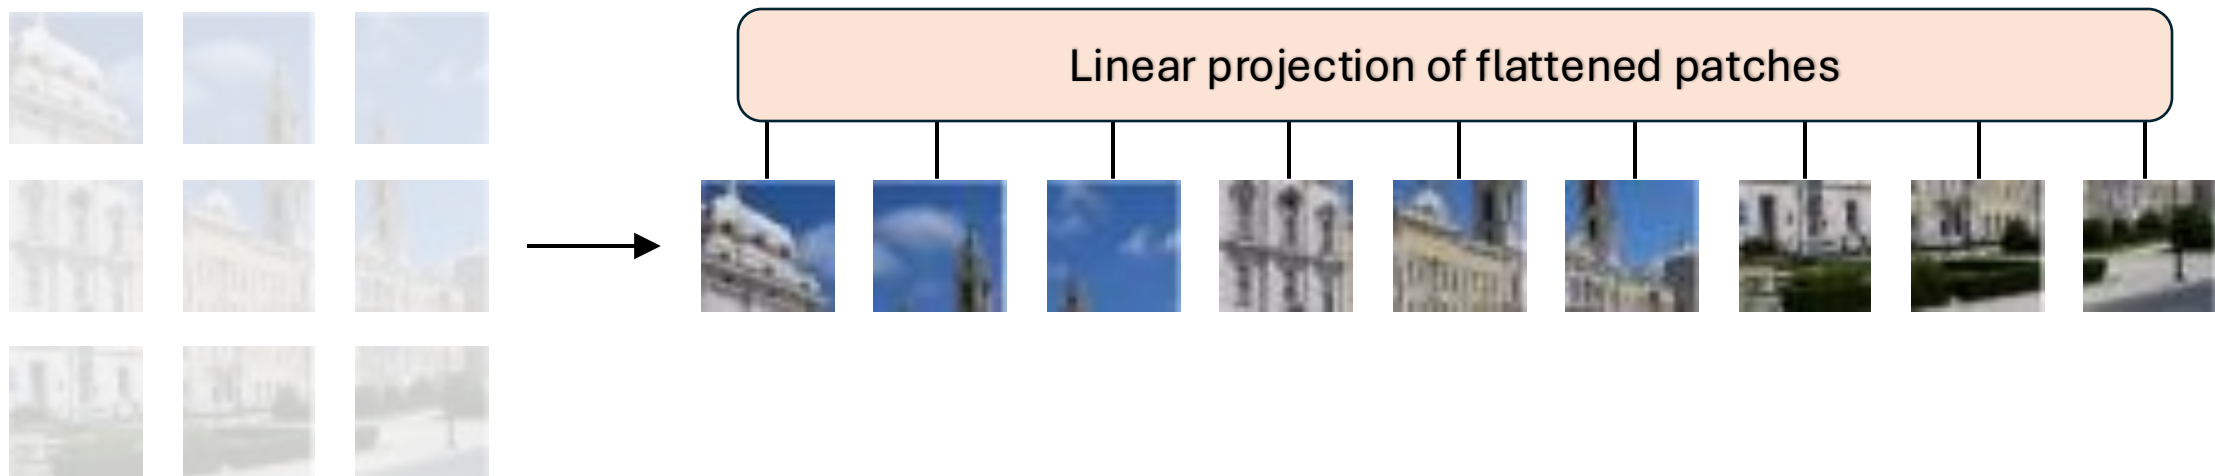
OVERVIEW OF THE ViT-ARCHITECTURE:



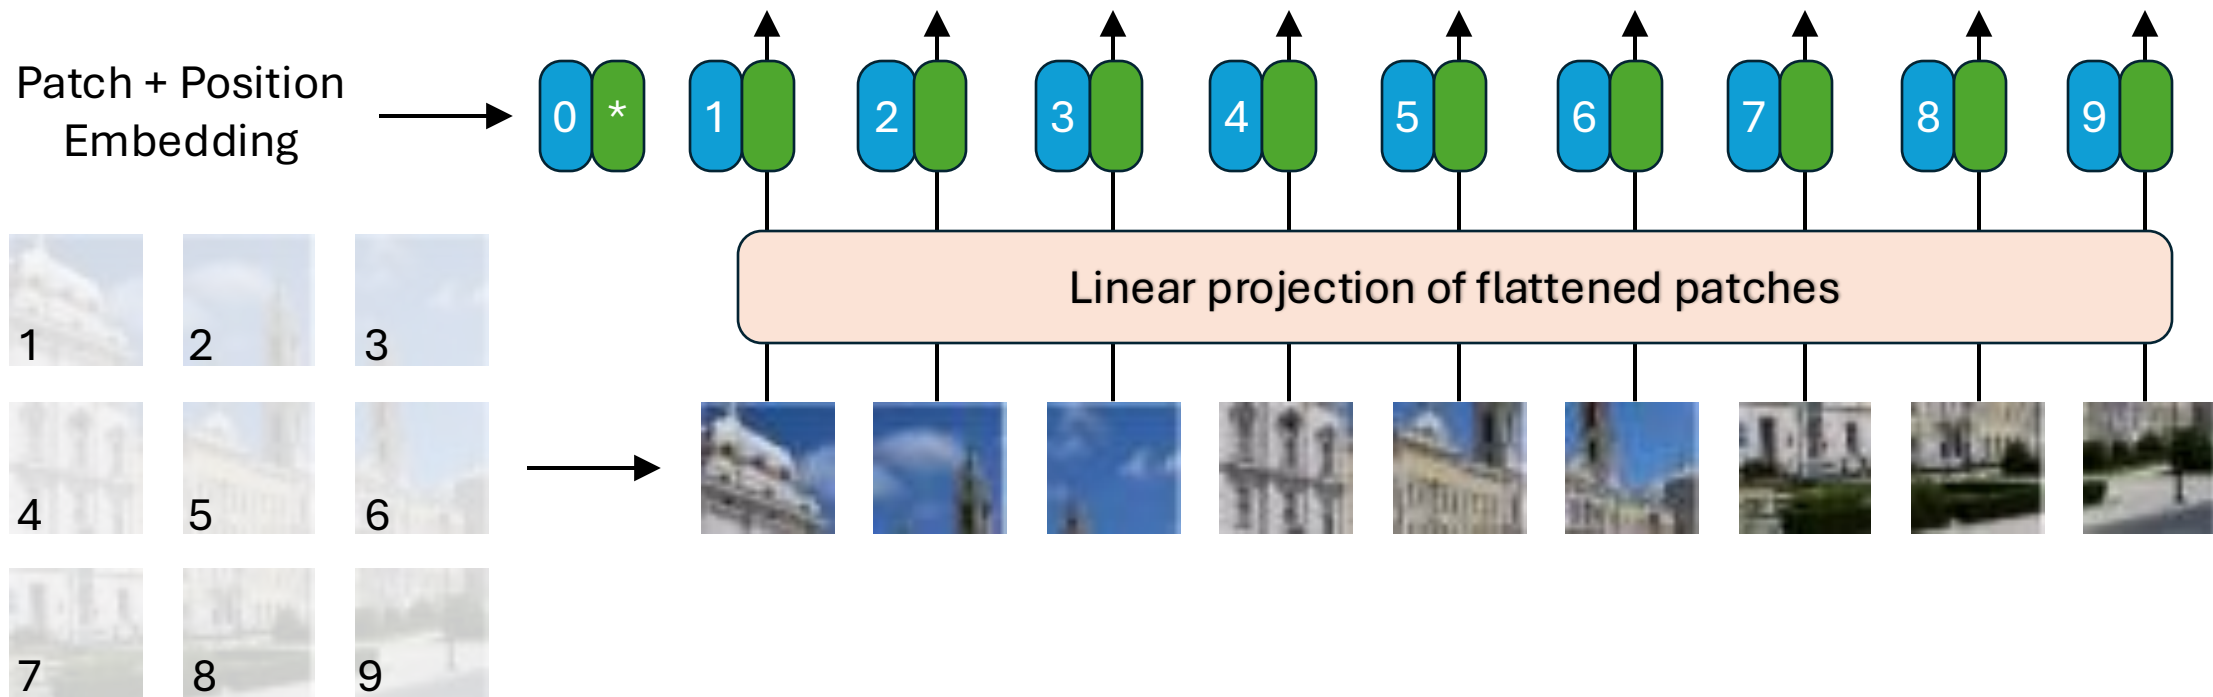
OVERVIEW OF THE ViT-ARCHITECTURE:



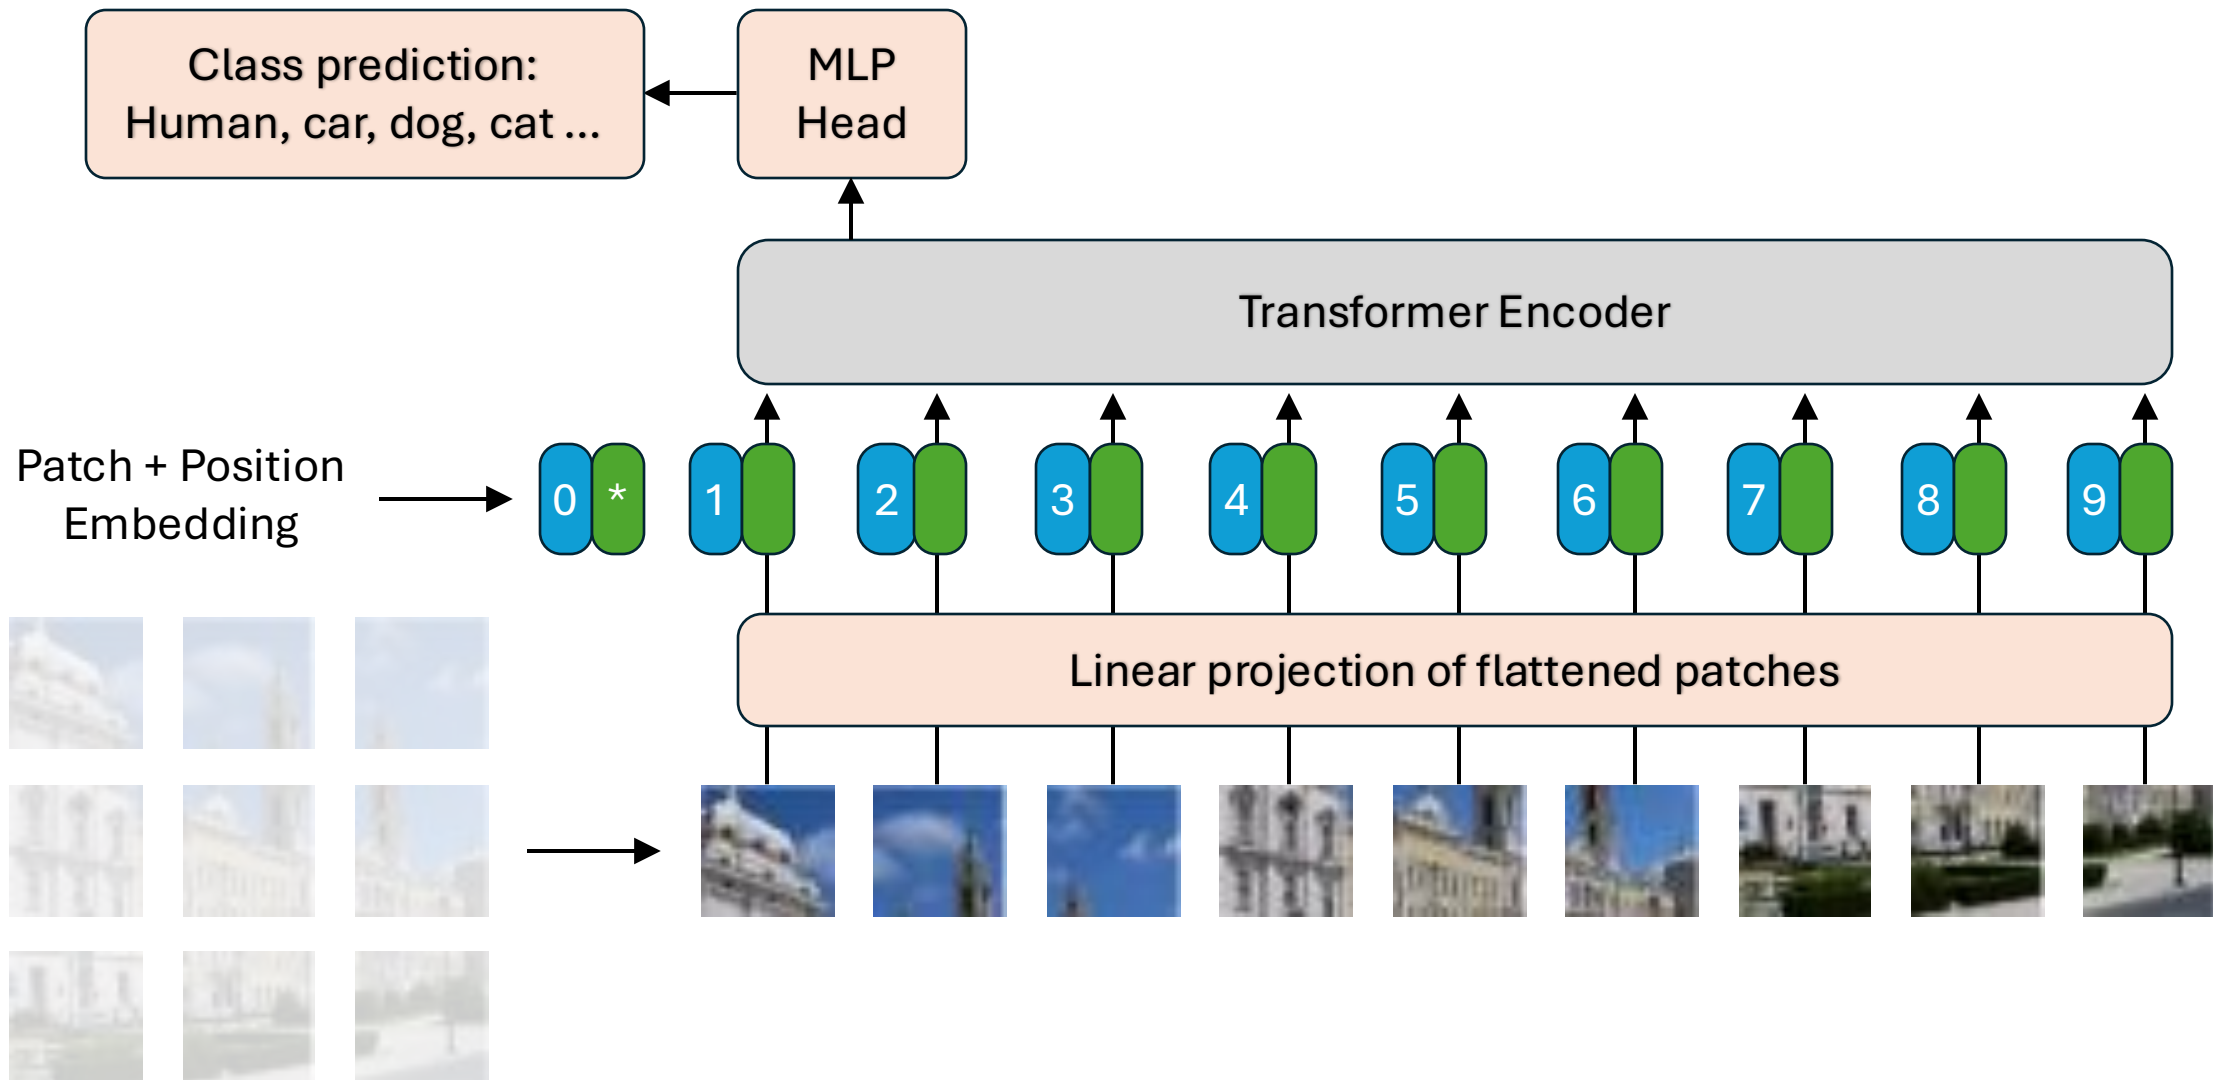
OVERVIEW OF THE ViT-ARCHITECTURE:



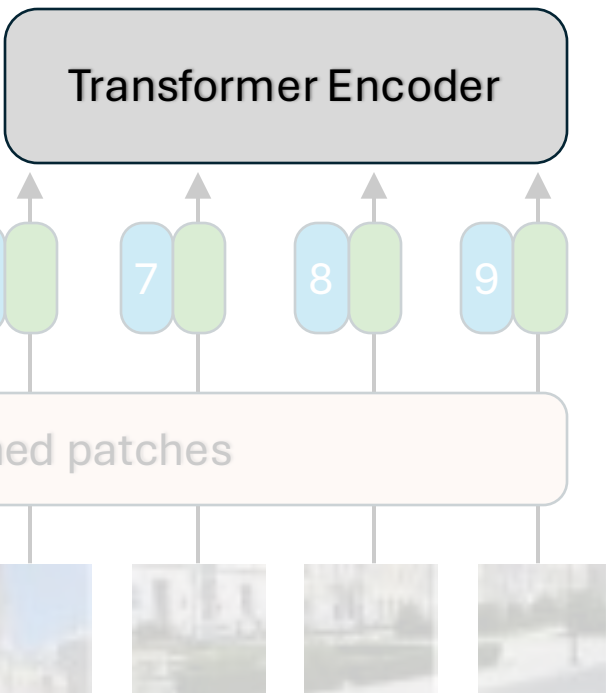
OVERVIEW OF THE ViT-ARCHITECTURE:



OVERVIEW OF THE ViT-ARCHITECTURE:



OVERVIEW OF THE ViT-ARCHITECTURE:



What is the “Transformer Encoder” how different / similar to NLP is it?

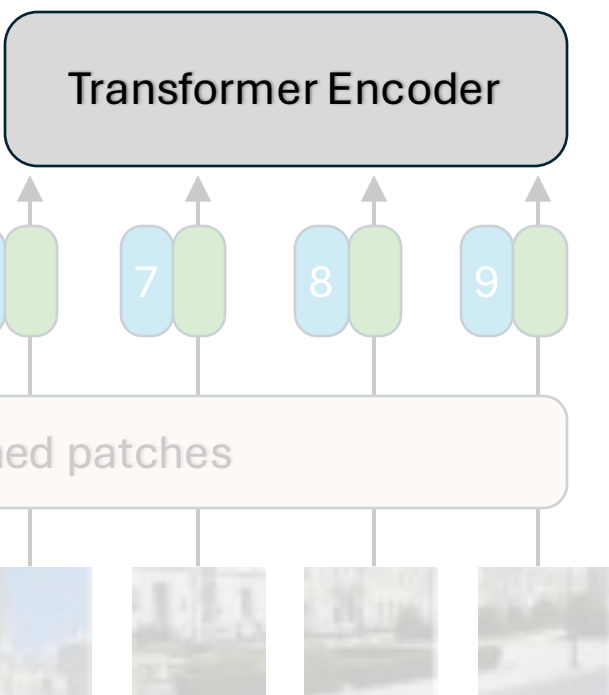
Both encoders are fundamentally similar in architecture and operation, but there are key differences due to the nature of the input data

Core Similarity

- **Multi-Head Self-Attention:** Computes relationships between all input tokens/patches, capturing global dependencies.
- **Feed-Forward Networks (FFNs):** Applied independently to each token/patch for feature transformation.
- **Layer Normalization and Residual Connections:** Stabilize training and improve gradient flow.
- **Positional Encoding/Embeddings:** Preserve sequence or spatial order.

The encoder processes a sequence of input tokens (or patches in ViTs) and produces contextualized representations by attending to all elements in the sequence simultaneously.

OVERVIEW OF THE ViT-ARCHITECTURE:



What is the “Transformer Encoder” how different / similar to NLP is it?

Differences

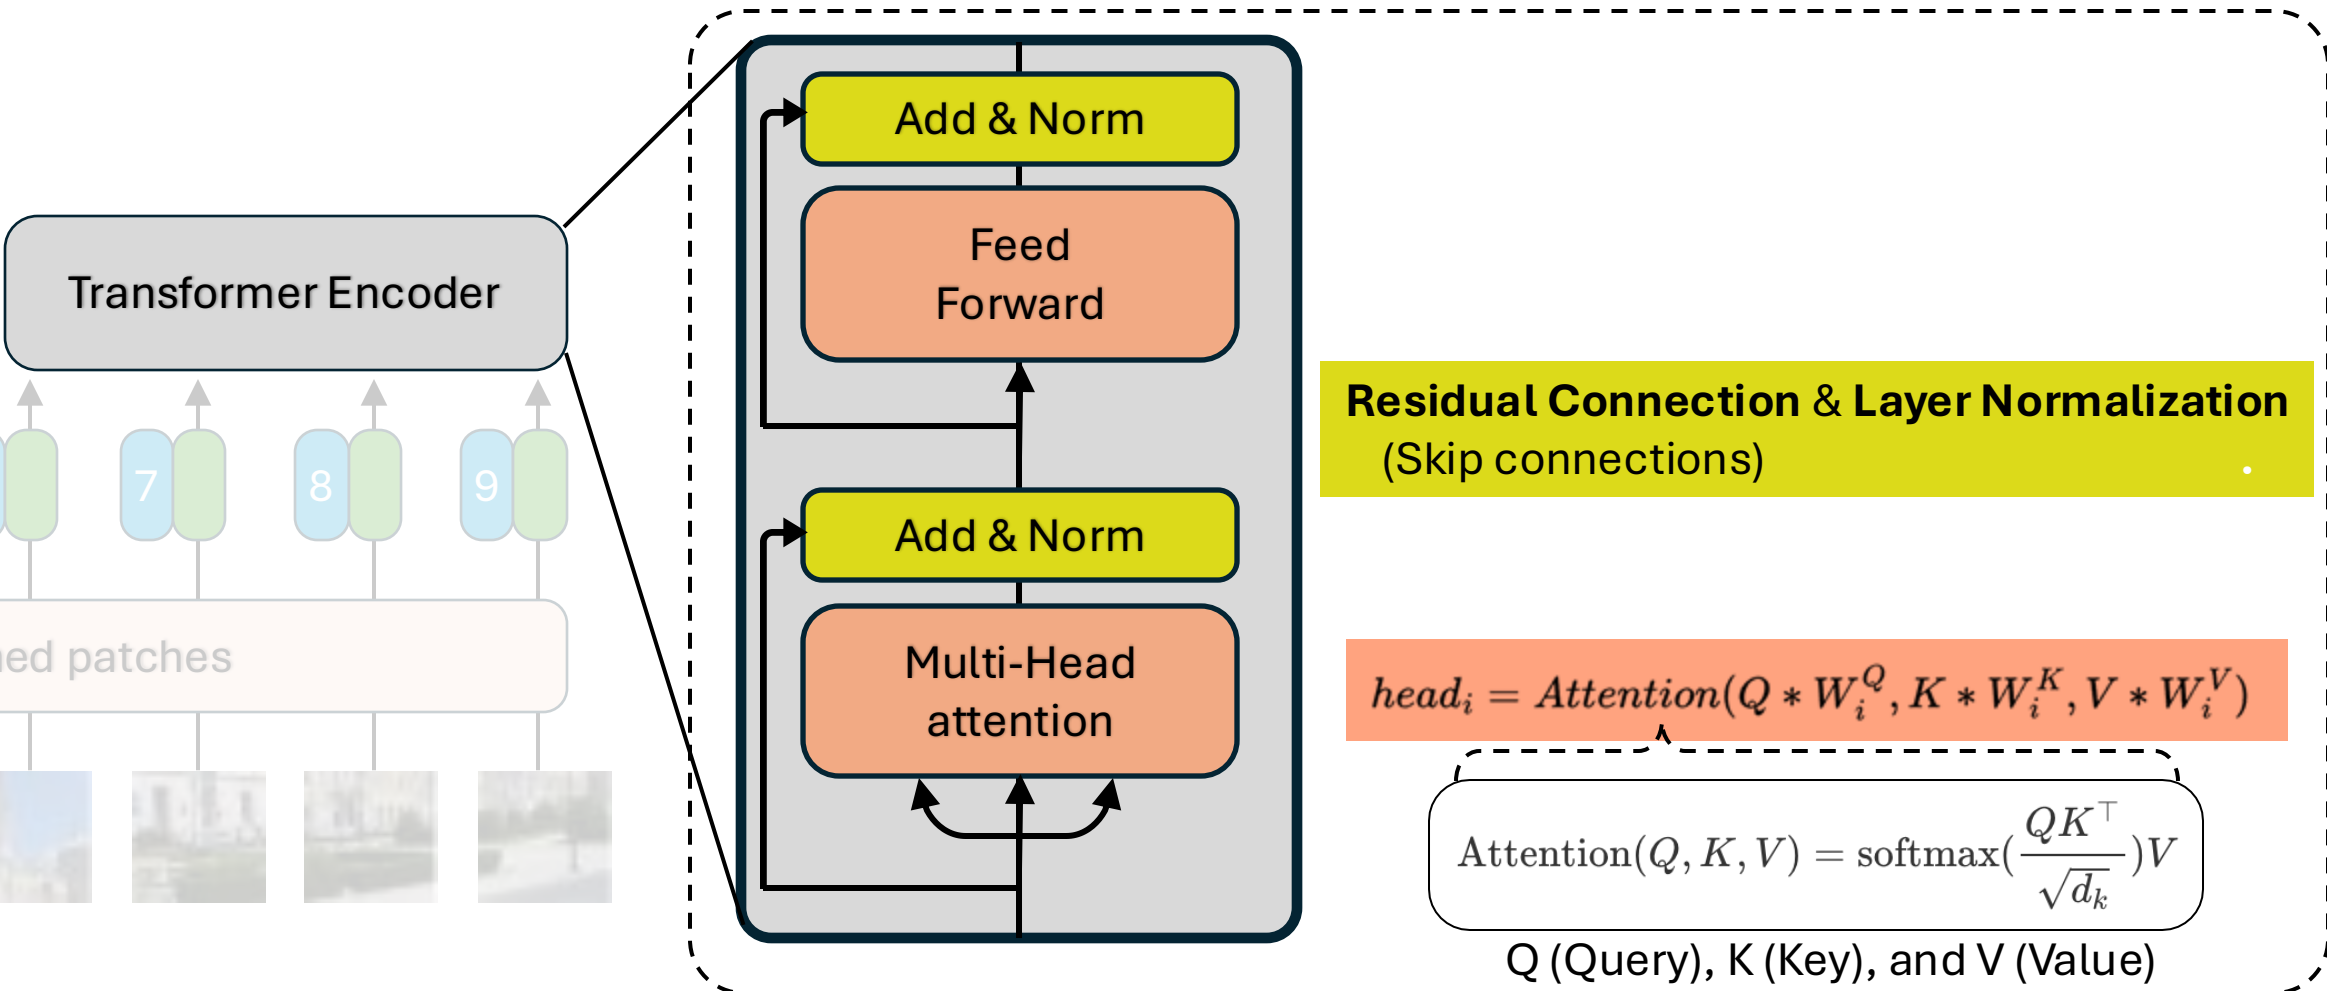
- Input

- **NLP:** Takes a sequence of tokenized text (such as words, subwords) and tokens are mapped to embeddings via a vocabulary lookup (e.g., WordPiece)
- **Vision:** Takes a sequence of image patches (e.g., 16x16 pixel patches). And the patches are flattened and linearly projected to a fixed-dimensional embedding space.

- Positional Encoding

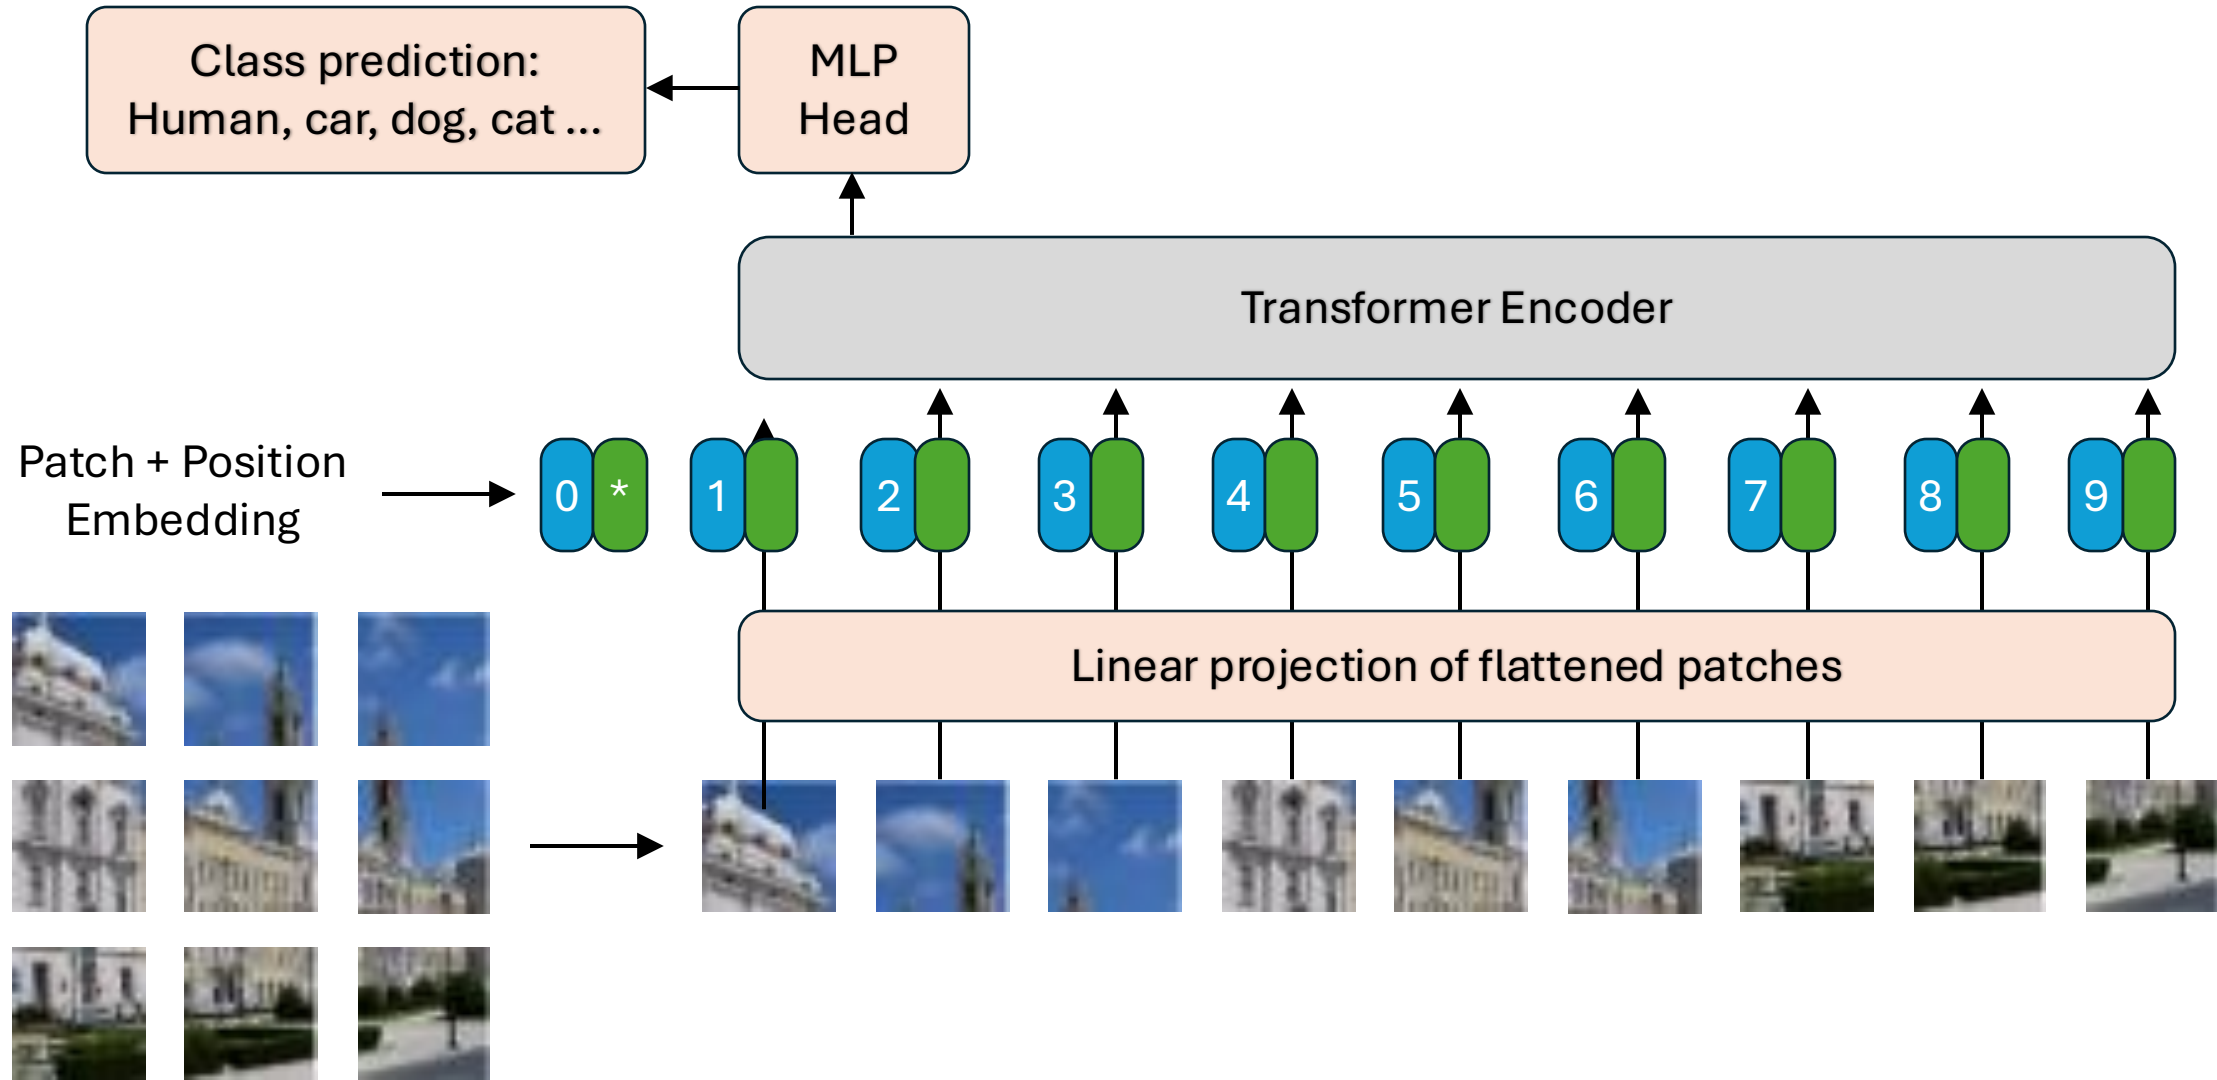
- **NLP:** In natural language, the input is a 1D sequence of words. The position can be represented by a single index, like 1, 2, 3, and so on. The positional encoding is therefore a 1D vector corresponding to each word's index.
- **Vision:** Images are 2D. The position of a patch is defined by its row and column coordinates. While some ViTs use a flattened 1D sequence of patches (e.g., from top-left to bottom-right), more advanced methods might use a 2D encoding that explicitly models the row and column positions

OVERVIEW OF THE ViT-ARCHITECTURE:




Divers ViTs architectures

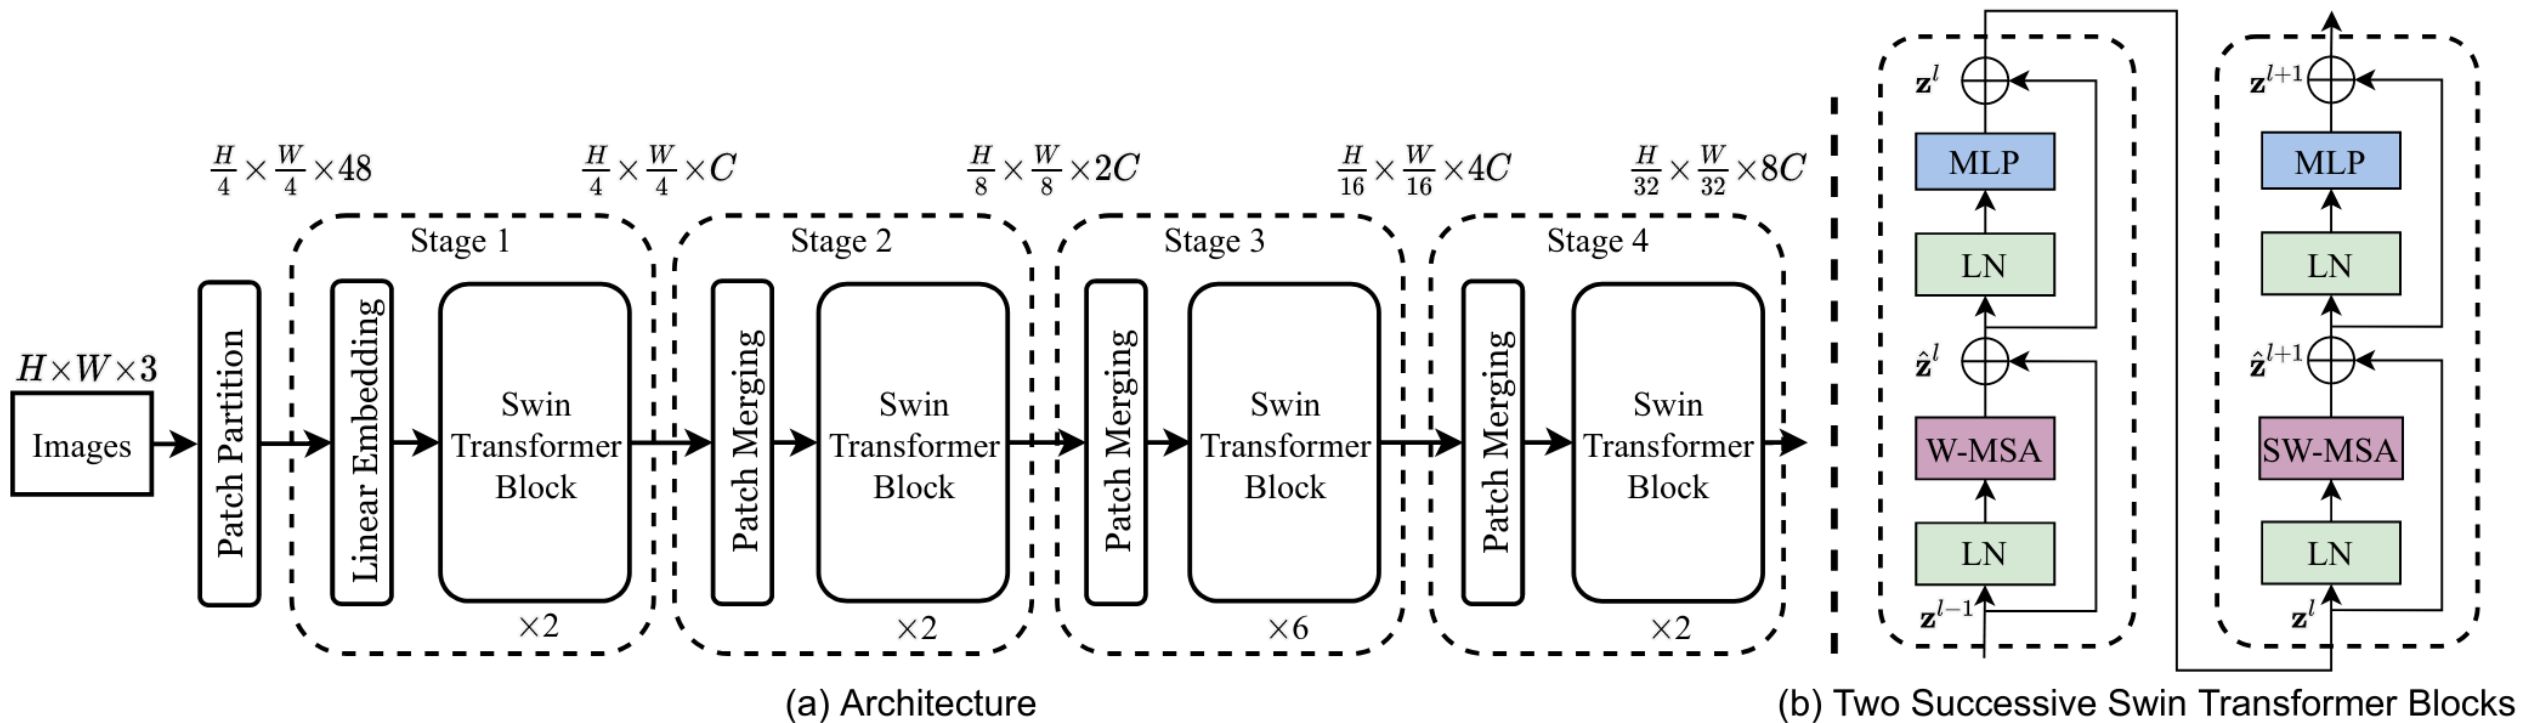
ViT : Google - Vision Transformer ([paper](#))



DeiT : Meta - Data efficient Image Transformer ([paper](#))



Microsoft – Swin-Transformer ([paper](#))





Hands-on:

(1) Pre-trained classification model





Hands-on:

(2) Visualization of features map





Hands-on:

(3) Finetuning ViT on classification task



ViTs in object detection


DETR: End-to-end Detection TRansformer ([paper](#))

Backbone

Encoder

Decoder

Prediction heads



DETR: End-to-end Detection TRansformer ([paper](#))


Backbone

Encoder

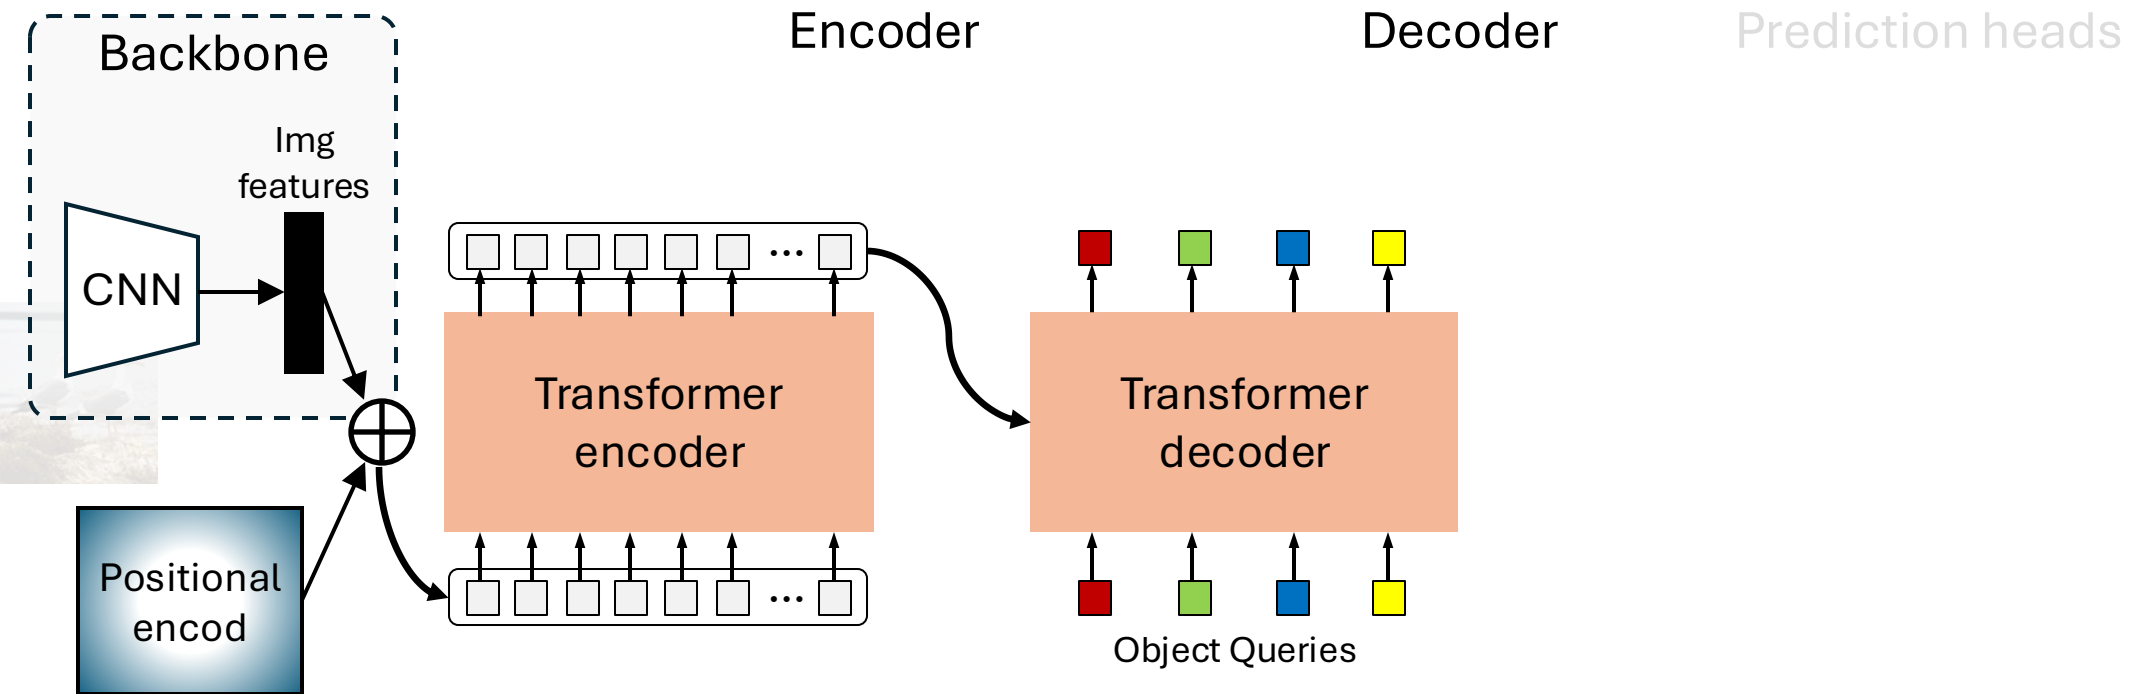
Decoder

Prediction heads

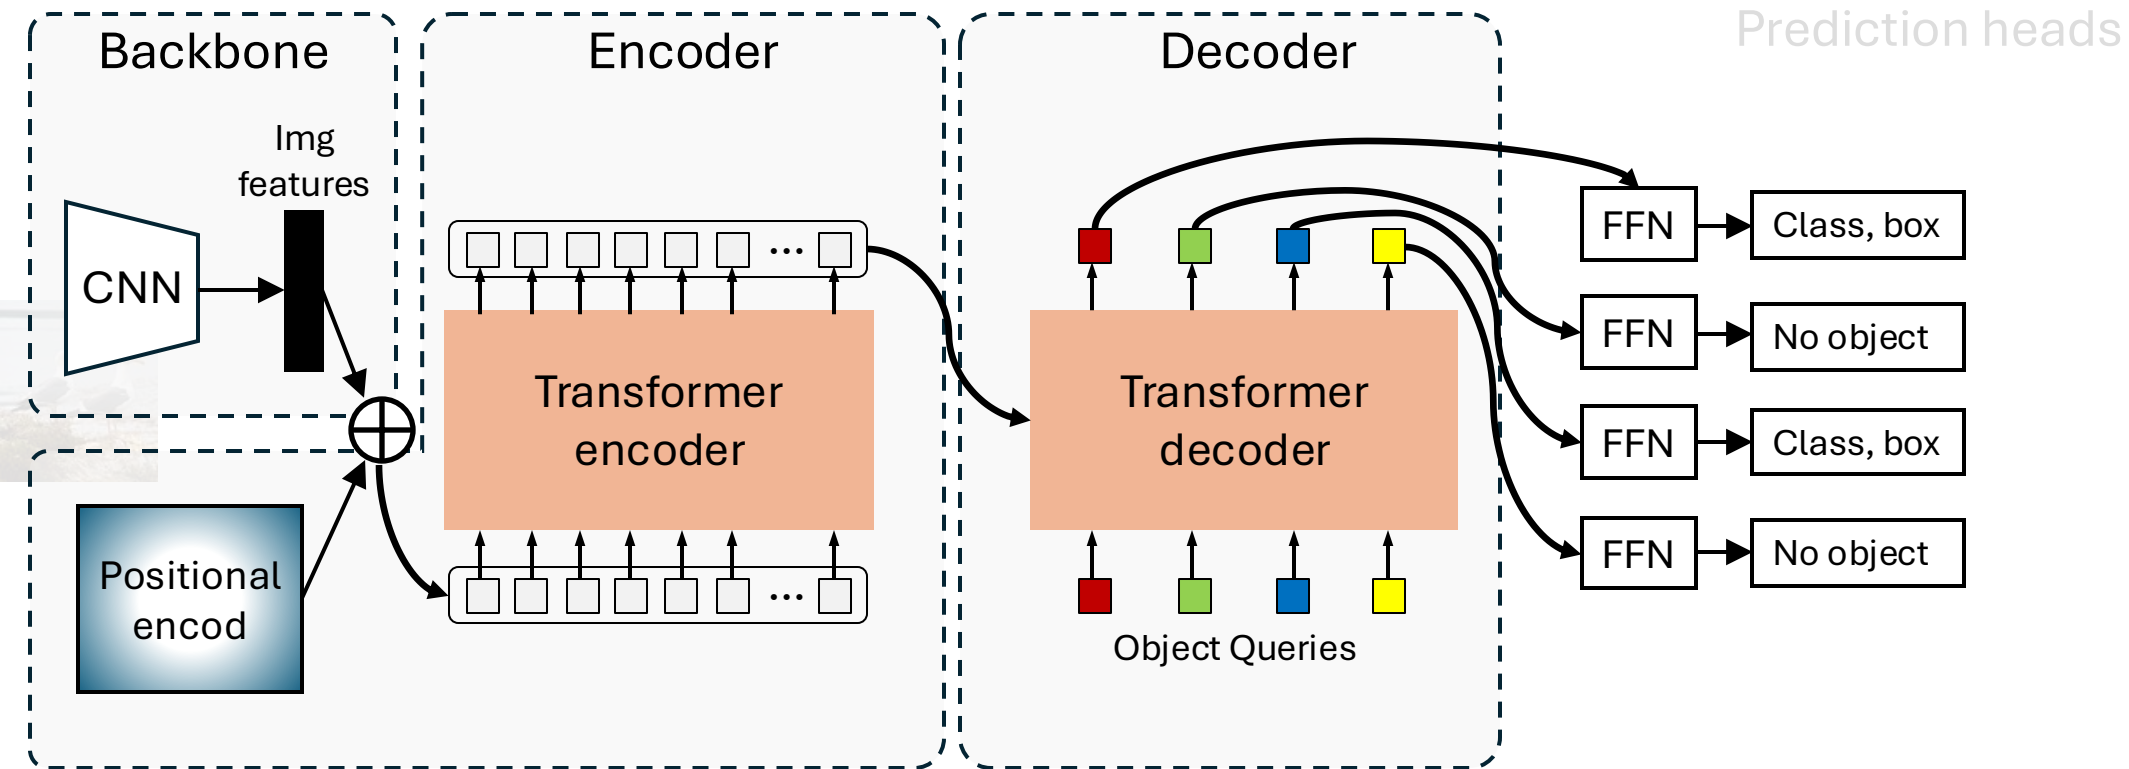
Pretrained CNN: VGG – ResNet ...etc



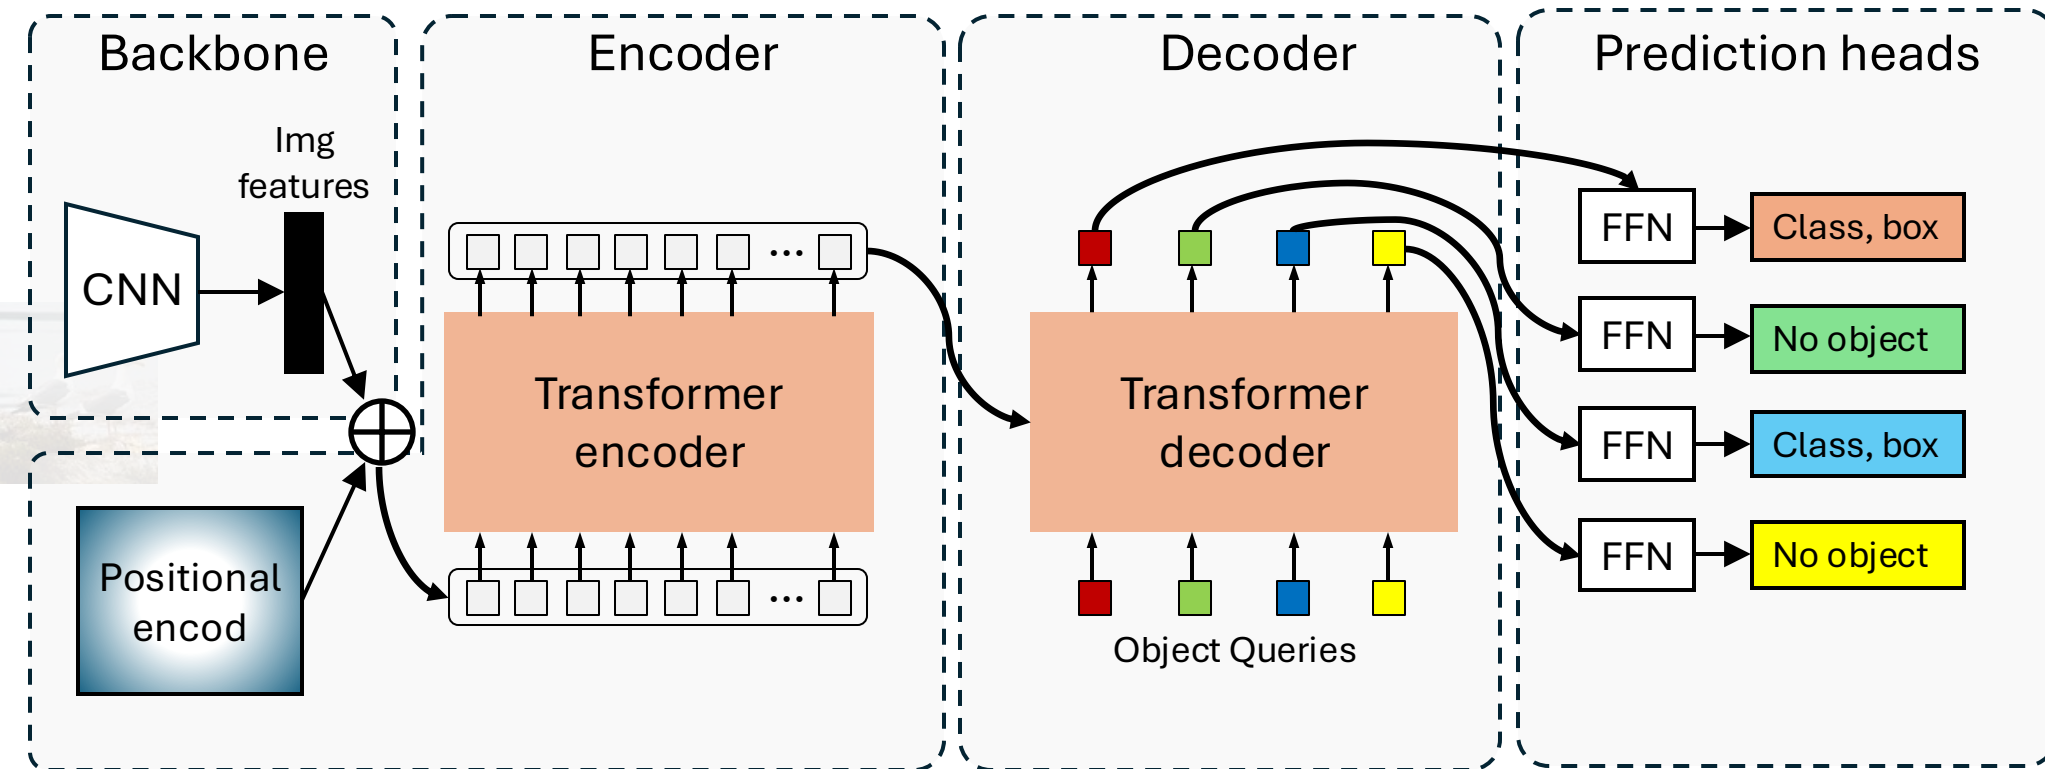
DETR: End-to-end Detection TRansformer ([paper](#))




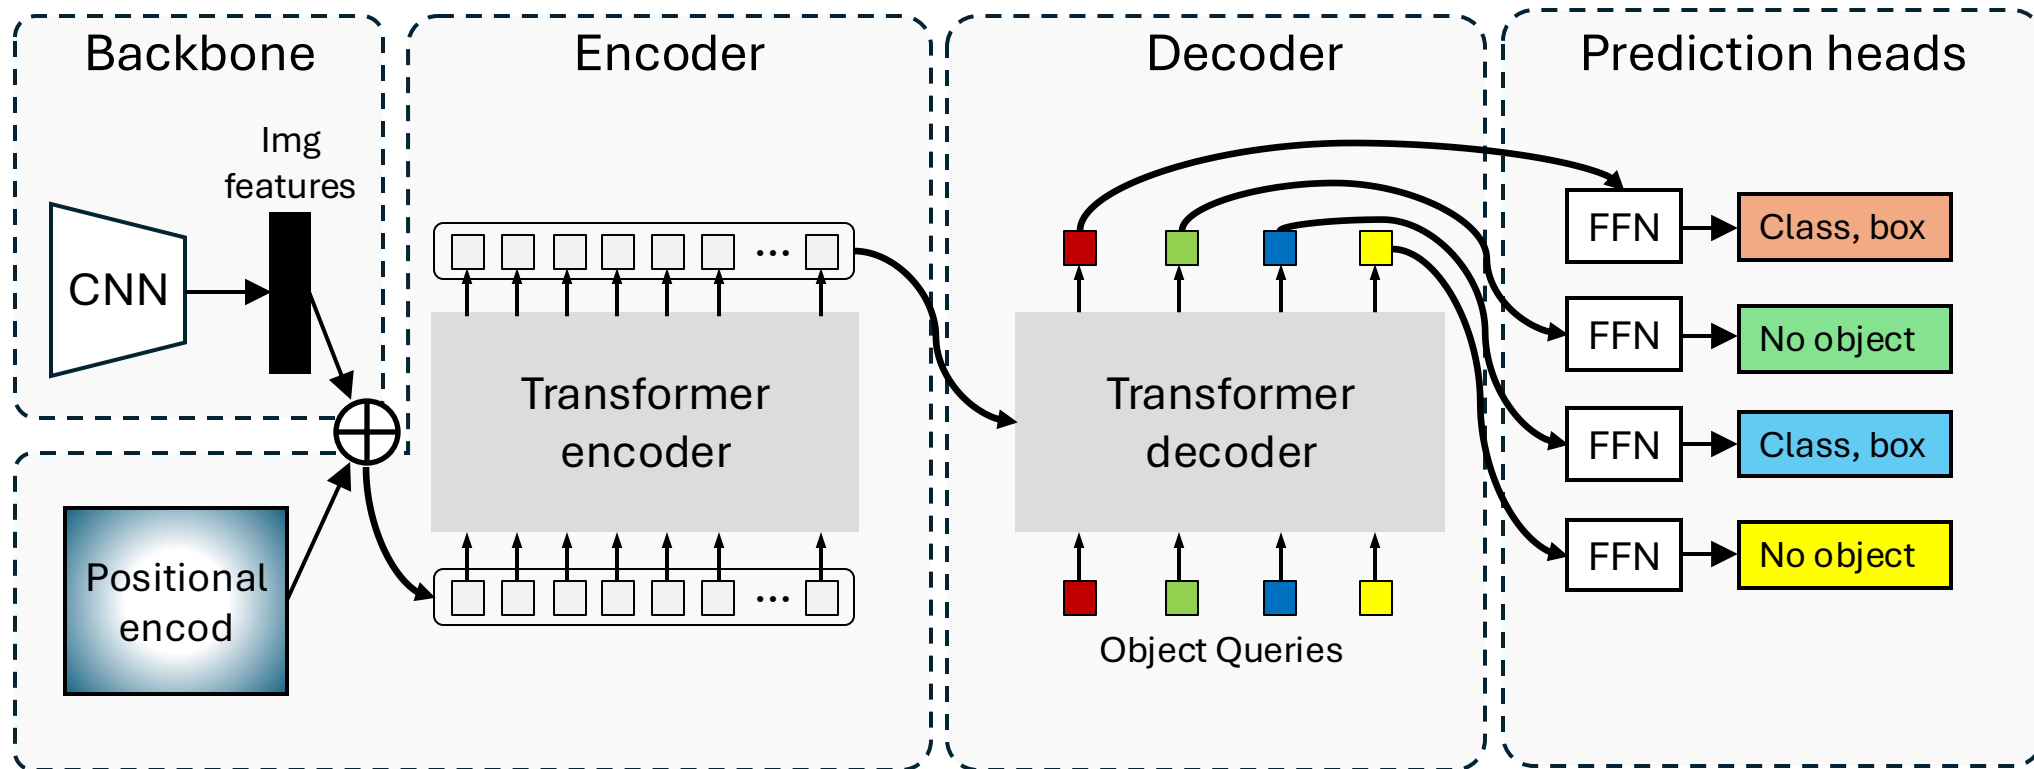
DETR: End-to-end Detection TRansformer ([paper](#))



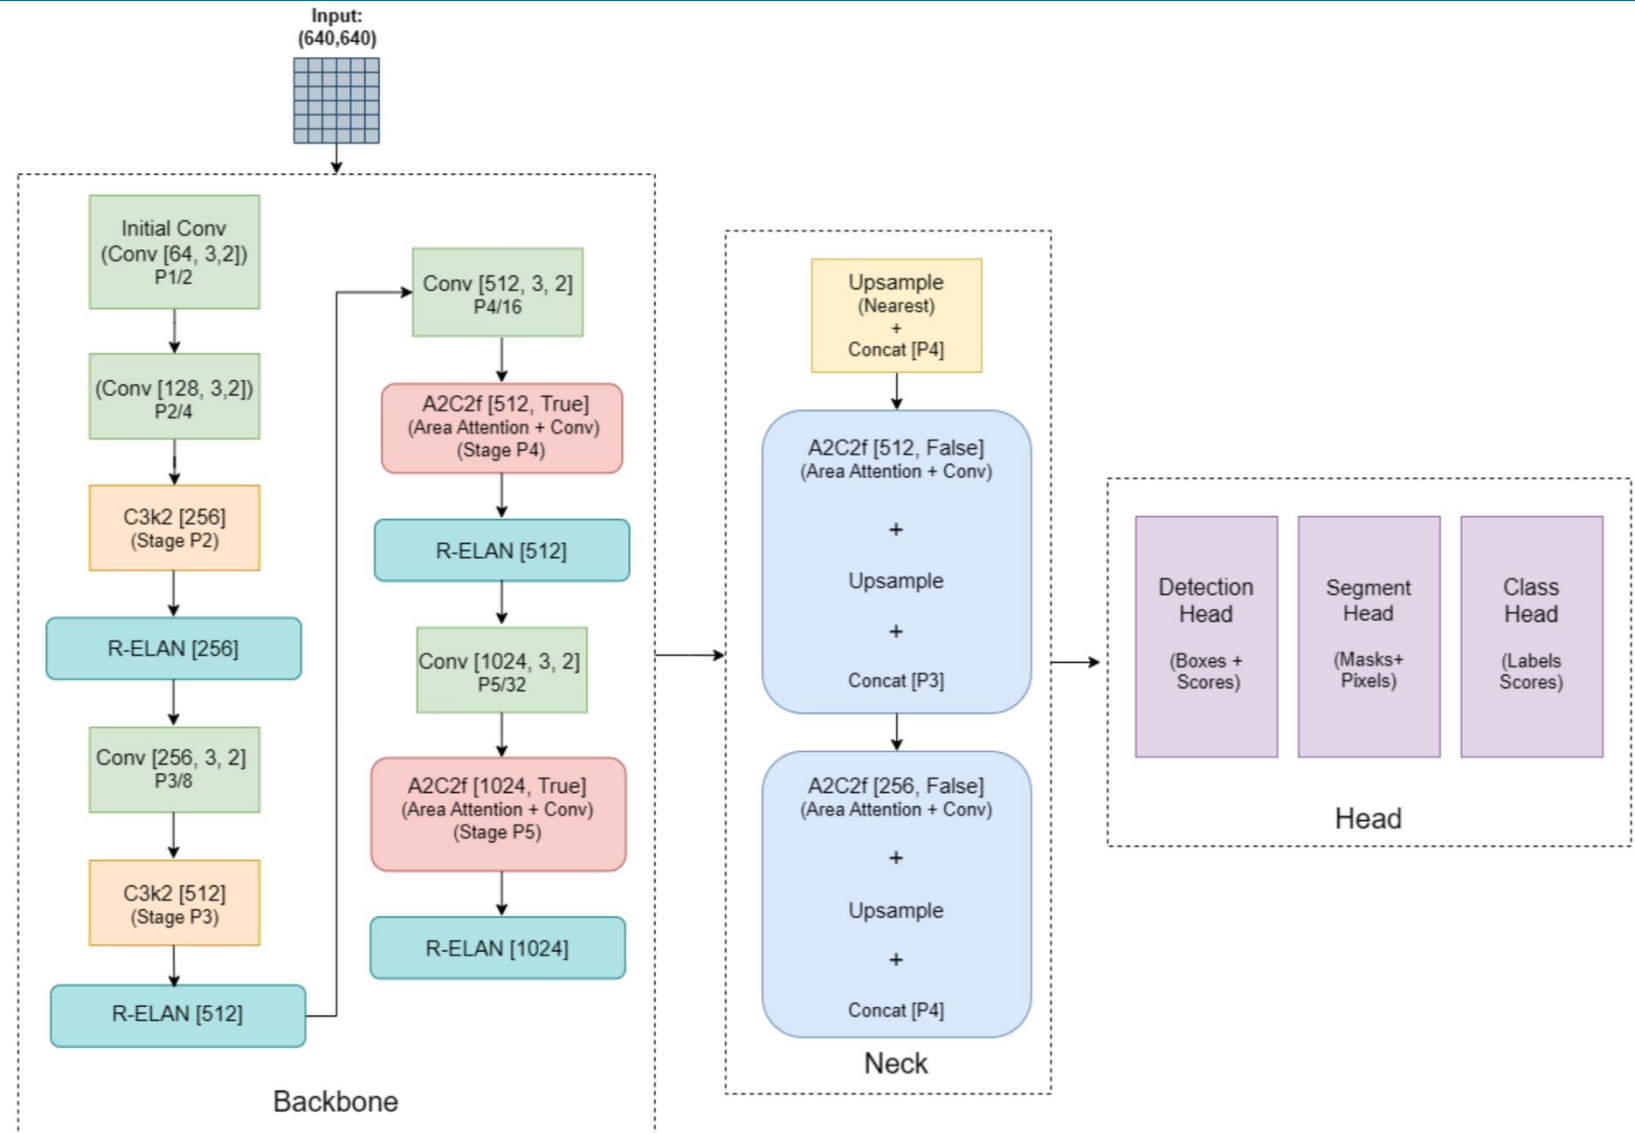
DETR: End-to-end Detection TRansformer ([paper](#))



DETR: End-to-end Detection TRansformer ([paper](#))



YOLO series - v12: Attention-centric object detection





Hands-on:

(1) Exploring models – inferencing



Hands-on:

(2) Finetuning object detection model on custom dataset



Hugging Face Transformers library


- Background & Motivations
- Applications: image captioning, retrieval, grounding, generative tasks
- CLIP and Zero-Shot Learning
- Hands-on:
 - Inference CLIP for zero-shot classification
 - GroundingDINO for zero-shot detection
 - Prompt Engineering
- BLIP and Multimodal Generation
- Diffusion Model – introduction to Markov Chain
- Hands-on:
 - Stable-diffusion
 - Prompt engineering

Background & Motivations



Introduction to Generative AI Fundamentals: Previous generative AI models rely on GAN models. Attention-Generative AI as models that create new data (e.g., images, text) from learned patterns in training data, contrasting it with discriminative models.

Motivations: solving data scarcity (e.g., synthetic data for training), creativity (art, design), personalization (e.g., custom content), and efficiency (automating tasks like content creation). Highlight real-world drivers like computational power growth, large datasets (e.g., LAION-5B), and ethical considerations (e.g., bias in generated outputs).

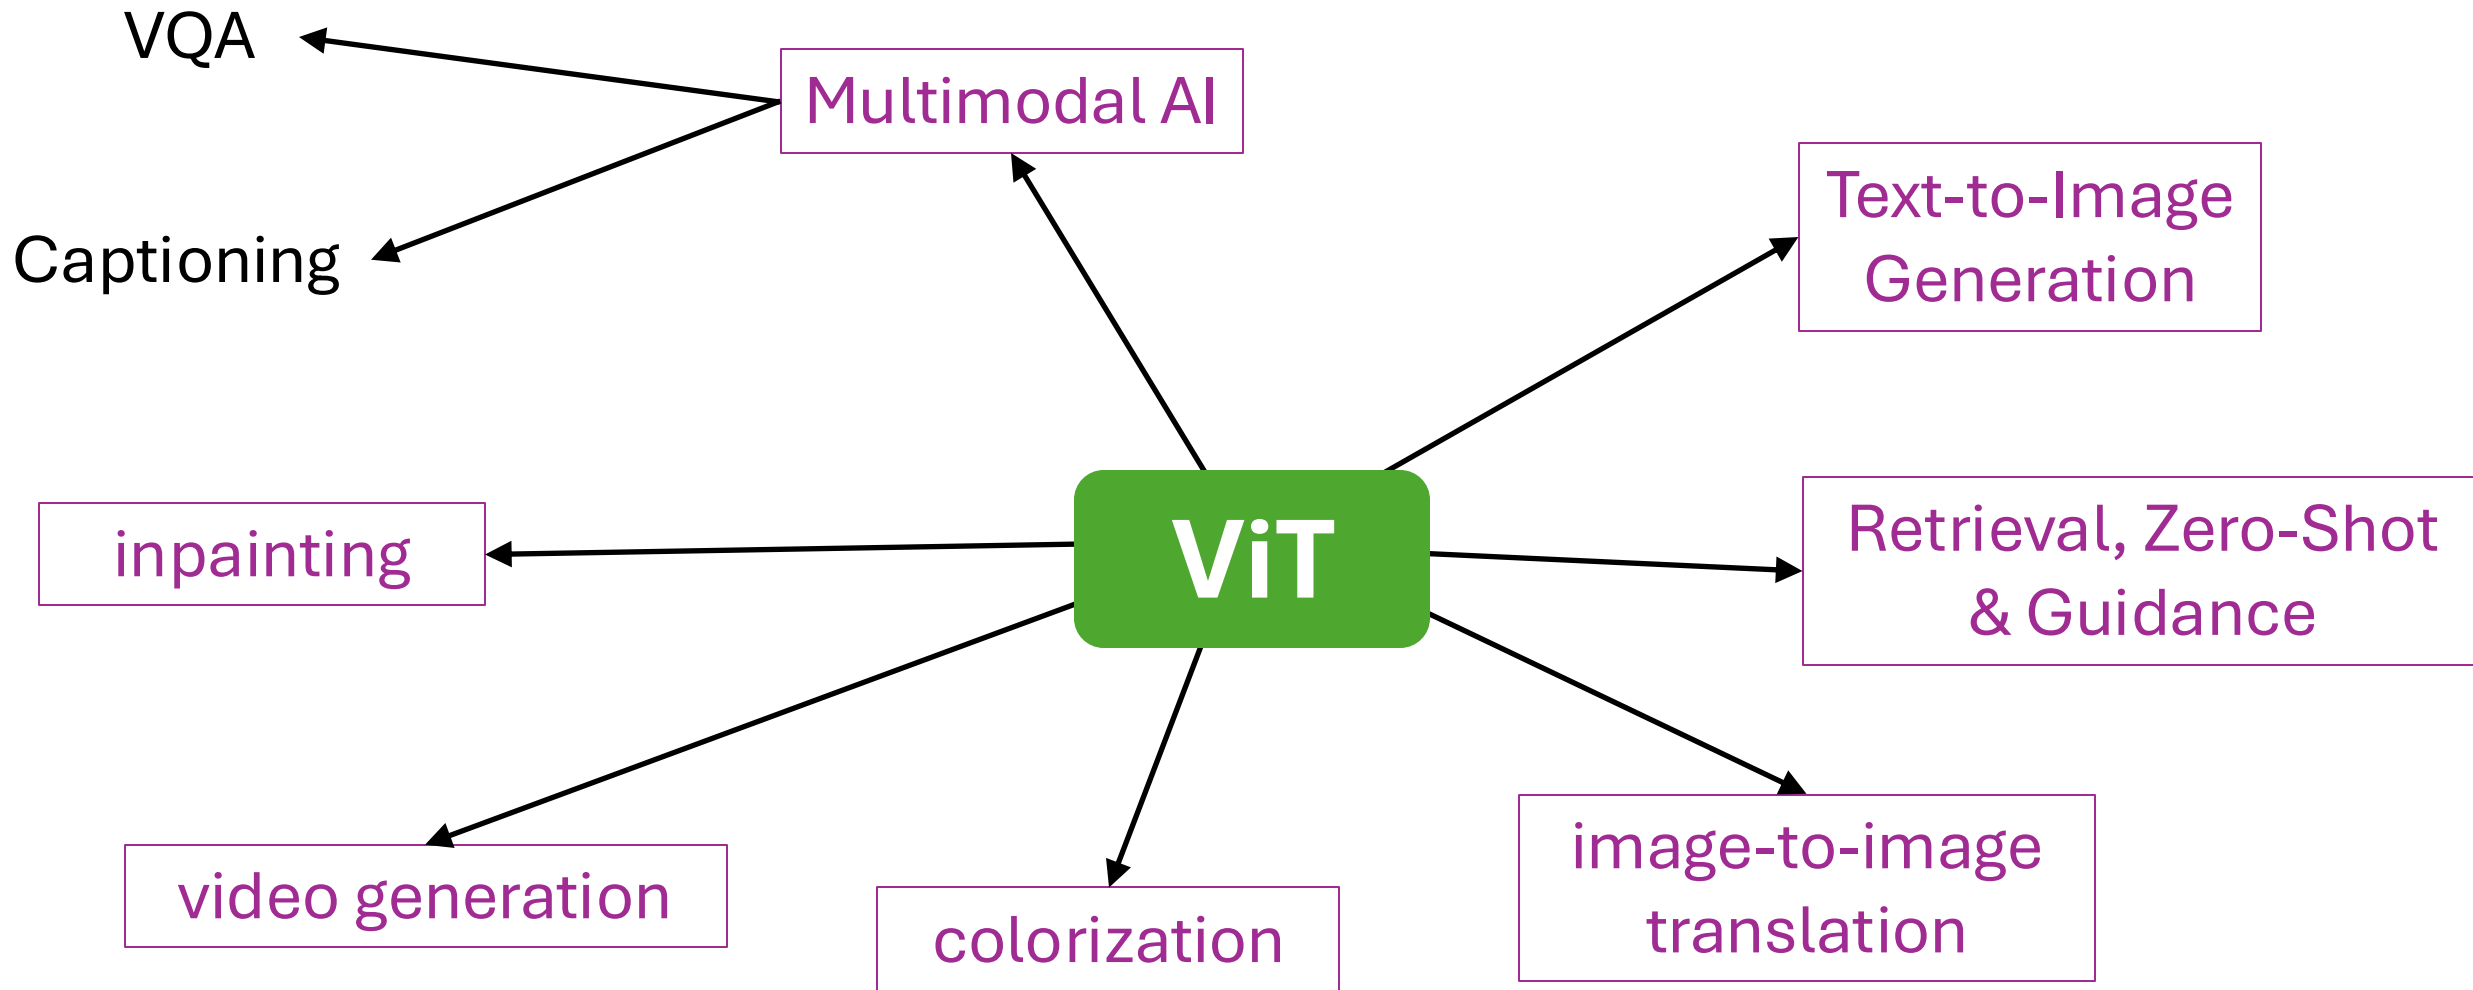
Synthetic image creation



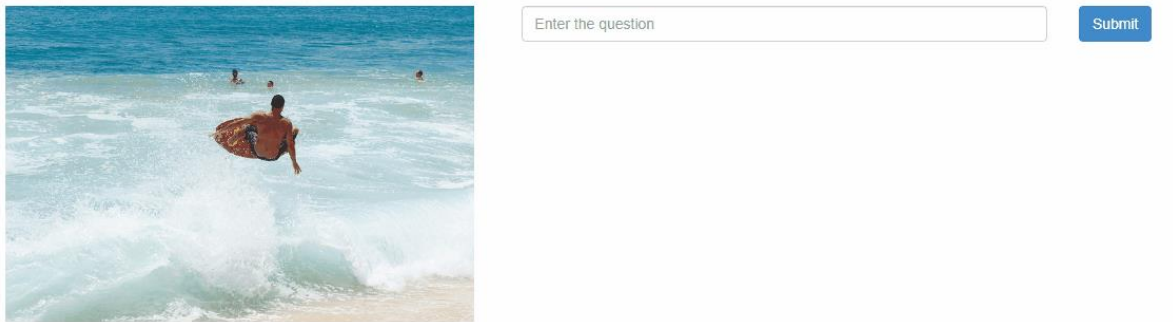
Captioning/VQA




Applications of Vision Transformers in Gen AI



Multimodal AI: VQA



inpainting



Text-to-Image Generation





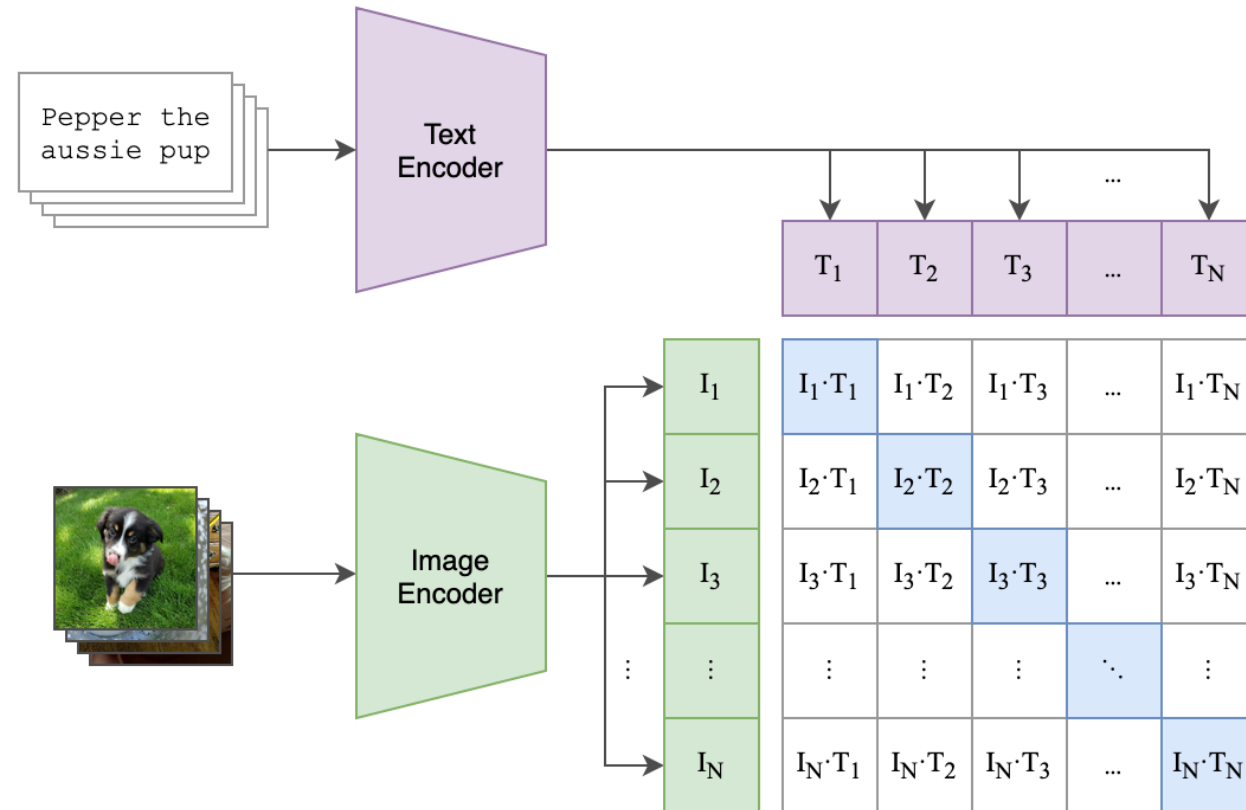
image-to-image translation



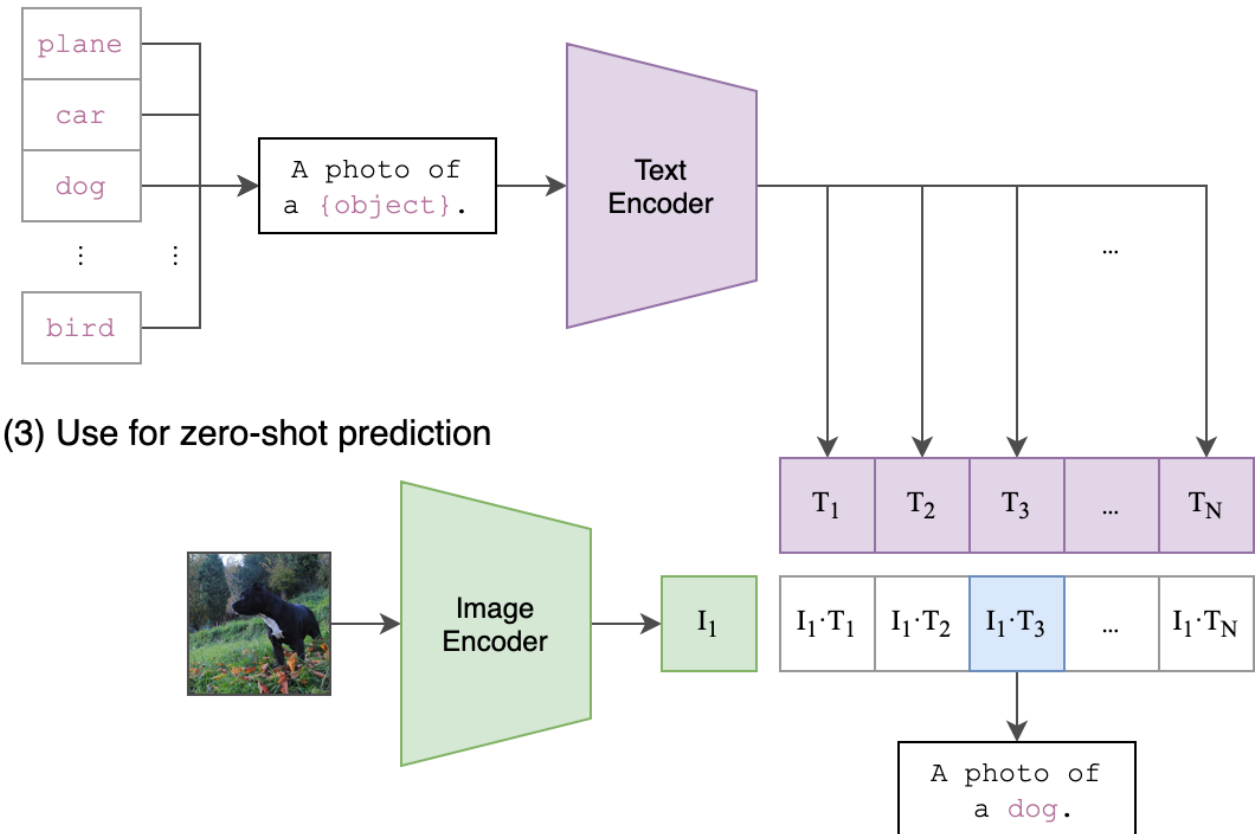
CLIP and zero-shot Learning

CLIP: Learning Transferable Visual Models From Natural Language Supervision ([Paper](#))


(1) Contrastive pre-training



(2) Create dataset classifier from label text

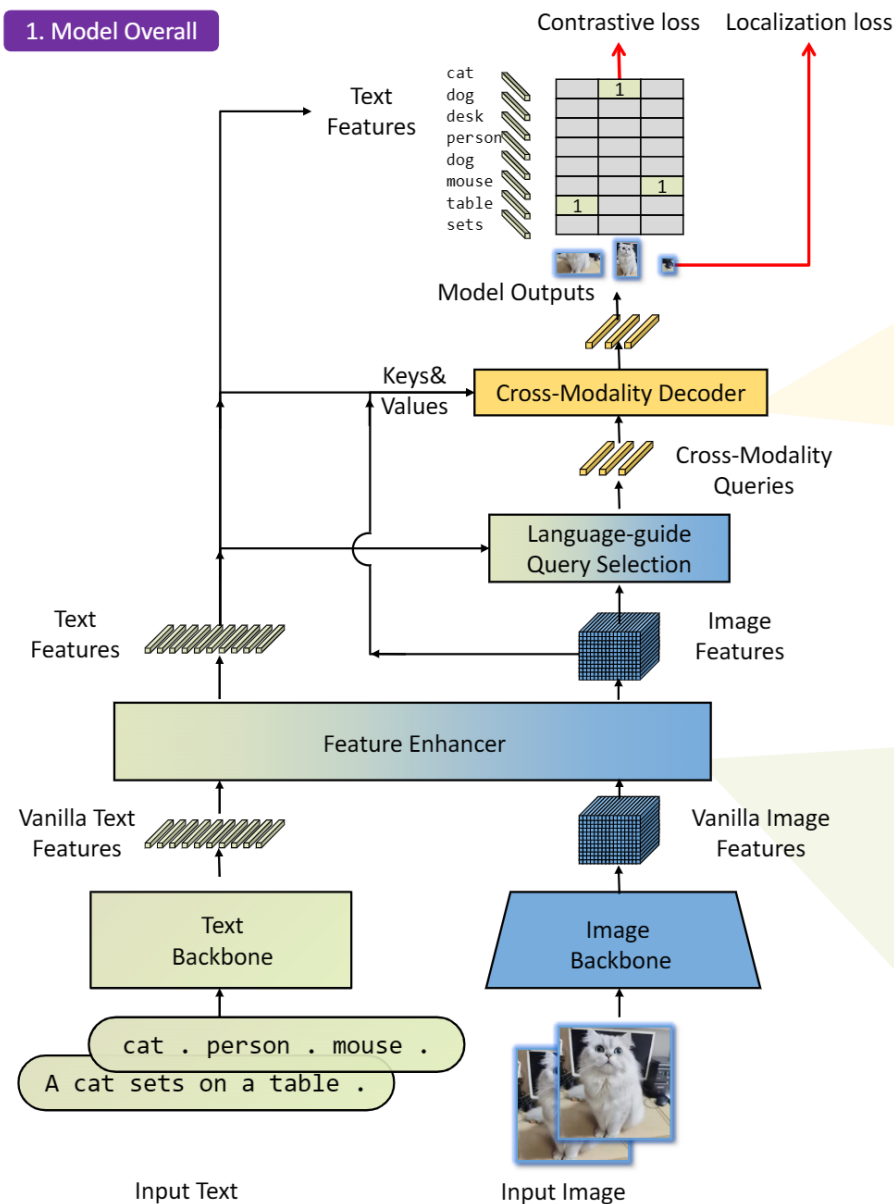


(3) Use for zero-shot prediction

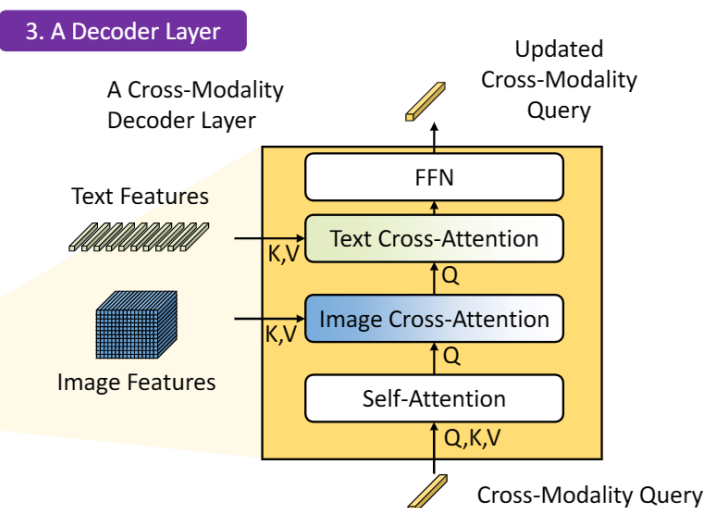


GroundingDINO: Marrying DINO with Grounded Pre-Training for Open-Set Object Detection ([Paper](#))

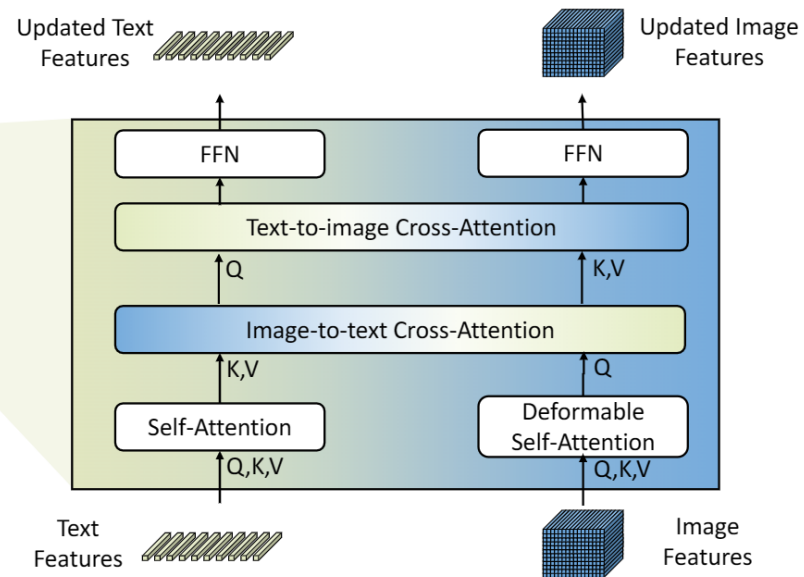
1. Model Overall



3. A Decoder Layer



2. A Feature Enhancer Layer





Hugging Face

Hands-on:

(1) CLIP for zero-shot classification



OpenAI



Hugging Face

Hands-on: (2) GroundingDINO for zero-shot detection



OpenAI



Hugging Face

Hands-on:

(3) Prompt Engineering

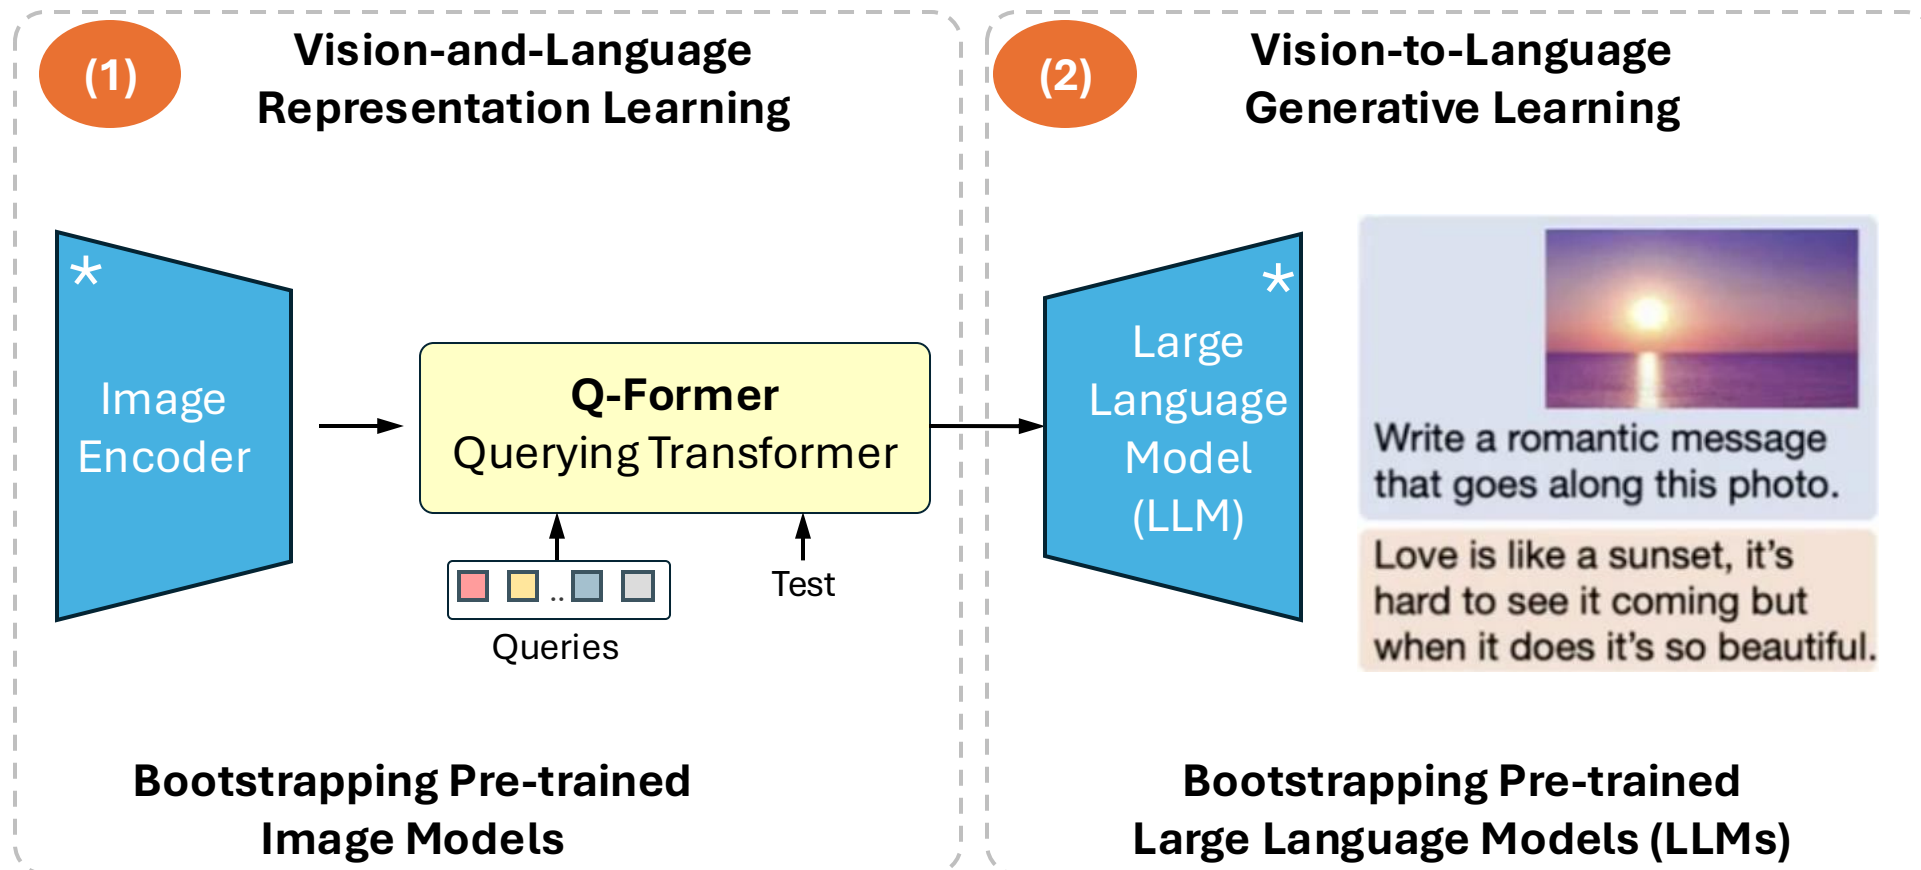


OpenAI

BLIP and multimodal generation

BLIP-2: Bootstrapping Language-Image Pre-training with Frozen Image Encoders and Large Language Models ([Paper](#))


Pre-train a light weight Q-Former in two stages:



	Concept	Paper
BLIP	A vision-language pretraining framework that uses a bootstrapped approach to generate and filter captions from noisy web data, enabling unified vision-language understanding and generation tasks like image captioning. It combines a vision encoder and text decoder for efficient captioning.	BLIP: Bootstrapping Language-Image Pre-training for Unified Vision-Language Understanding and Generation
IDEFICS	A multimodal model designed for vision and language tasks, including image captioning, built by fine-tuning open-access models like LLaVA and CLIP-ViT. It processes images and text jointly, leveraging a large-scale pretraining dataset for robust performance.	IDEFICS: An Open-Source Visual Language Model
GIT	A generative image-to-text transformer that excels in image captioning by directly generating text from image inputs. It uses a single vision transformer and a text decoder, optimized for lightweight and efficient captioning with performance comparable to BLIP-2.	GIT: A Generative Image-to-text Transformer for Vision and Language
InstructBLIP	An enhanced version of BLIP-2, fine-tuned with instruction-tuning datasets to improve zero-shot and few-shot performance in image captioning and other vision-language tasks. It leverages BLIP-2's architecture with additional task-specific optimizations.	InstructBLIP: Towards General-purpose Vision-Language Models with Instruction Tuning
COCA	A contrastive captioning model that combines contrastive learning with generative captioning. It uses image-text alignment to produce high-quality captions, offering performance close to BLIP-2 with lower computational requirements.	CoCa: Contrastive Captioners are Image-Text Foundation Models

Diffusion Models

Diffusion Models work by destroying training data through the successive addition of Gaussian noise, and then learning to recover the data by reversing this noising process.



Forward diffusion process

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t \mathbf{I})$$

distribution of
the noised
images


Output

Mean
 μ_t

Variance
 Σ_t

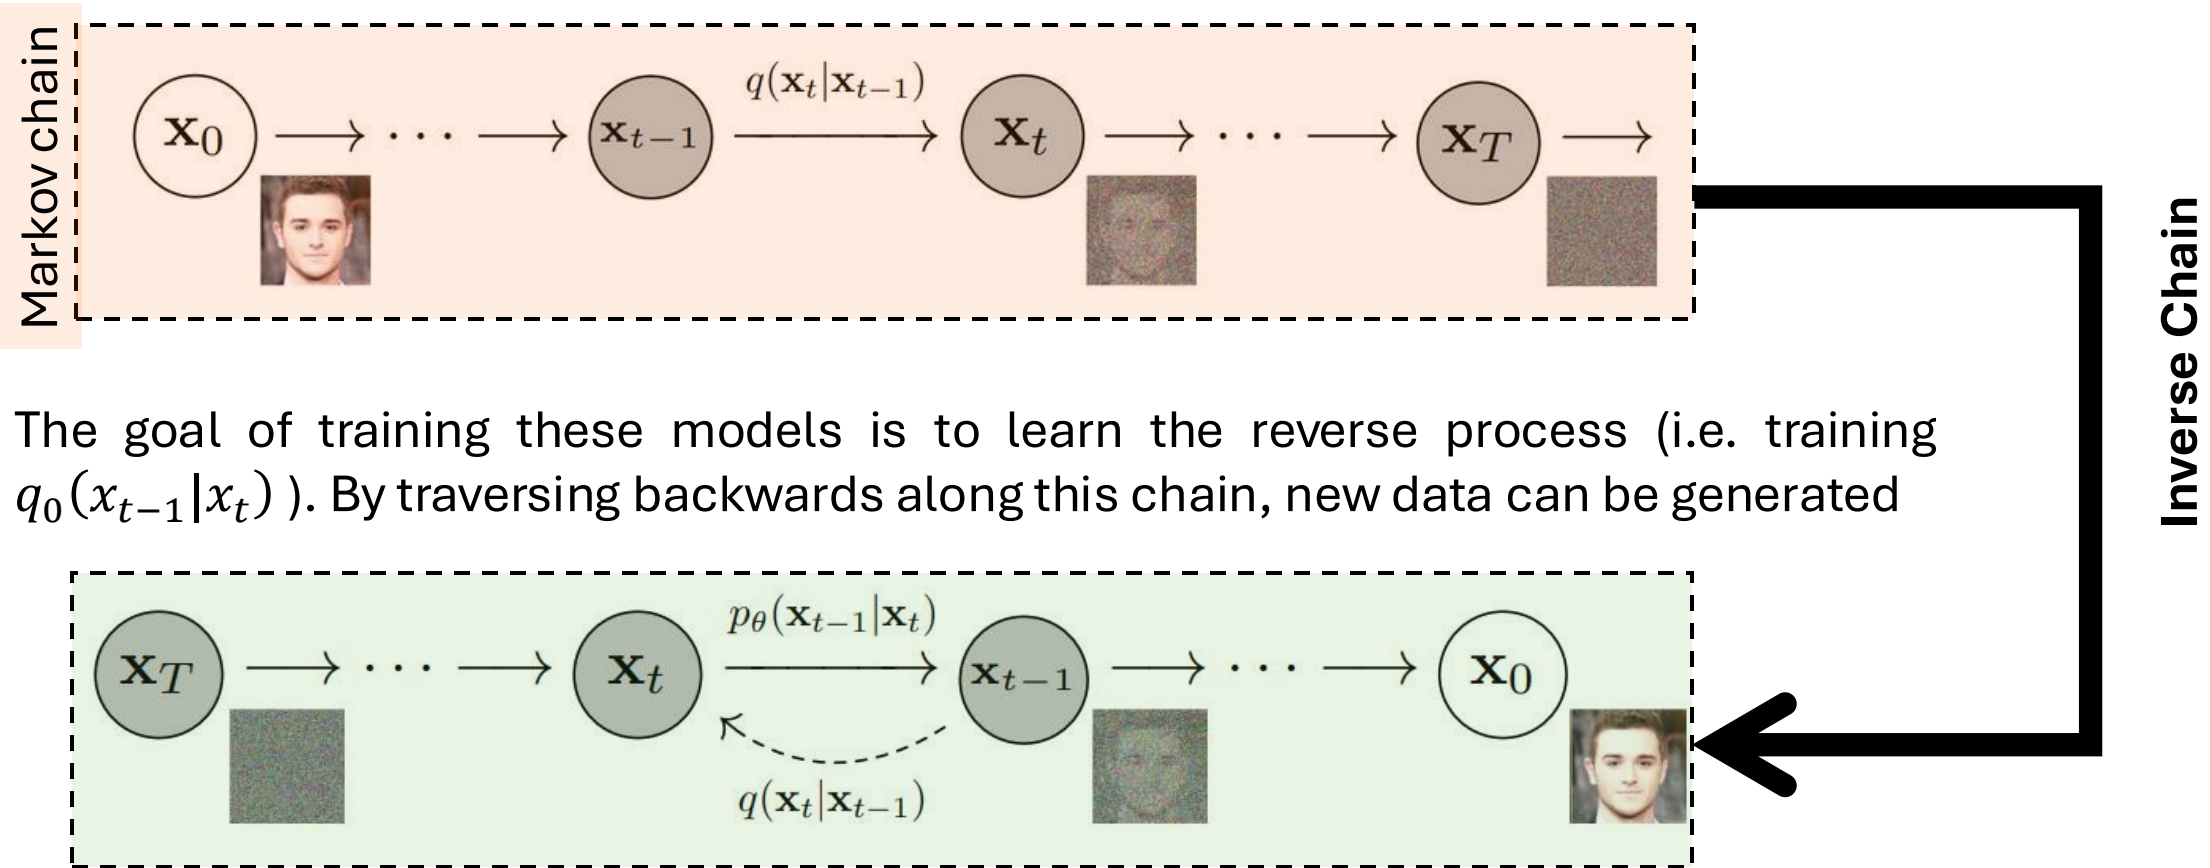
- t : This represents the **time step**. It goes from a clear image at $t=0$ to a completely noisy image at $t=T$.
- x_0 : This is your **original, clean data sample**. It's the starting point of the whole process.
- β_t : This is the **variance schedule**. It's a list of small numbers that control how much noise is added at each step. (1) It's a fixed sequence (not learned). (2) it typically starts with a very small number ($\beta_0 \approx 0$) so the first few steps add very little noise. (3) It gradually increases to a larger number ($\beta_T \approx 1$), so the later steps add a lot of noise. (4) The values of β_t are always between 0 and 1.
- \mathbf{I} : This is the **identity matrix**. It ensures that the noise is added independently to each pixel of the image, without any correlation between pixels.

Diffusion Models work by destroying training data through the successive addition of Gaussian noise, and then learning to recover the data by reversing this noising process.



The goal of training these models is to learn the reverse process (i.e. training $q_0(x_{t-1}|x_t)$). By traversing backwards along this chain, new data can be generated

Diffusion Models work by destroying training data through the successive addition of Gaussian noise, and then learning to recover the data by reversing this noising process.





Hugging Face

Hands-on:

(1) BLIP image captioning





Hugging Face

Hands-on:

(2) Stable diffusion



OpenAI

Computer Vision – Fundamentals

Introduction To Transformers In Vision

Program directors:



Dr. Prof.
Umberto Michelucci



Dr.
Aygul Zagidullina



Dr.
Safouane El Ghazouali

